



Jose Parreño Garcia <joparga3@gmail.com>

[TEST qpa - Paid] OpenAI and Anthropic are becoming consultancies (and you shouldn't be worried about it)

1 mensaje

Jose Parreño Garcia from Senior Data Science Lead <joseparreogarcia@substack.com>
Para: joparga3@gmail.com

18 de junio de 2026 a las 20:01

[View in browser](#)

The Senior Data Science Lead newsletter

Hi! It's Jose from Senior Data Science Lead newsletter, where I write about senior data science leadership, sharing hard-earned lessons from building teams, systems, and real-world machine learning.

Because you are a paid subscriber, you are directly supporting this work — and you get full access to the practical tools built alongside the writing:

- ***The full Data Storytelling course, focused on turning analysis into insights that land with executives. Because today, anyone can build a chart with AI, but very few people can build the right chart.***



Master Data Storytelling

JOSE PARREÑO GARCIA · NOVEMBER 27, 2025

[Read full story](#)

- ***50+ Notion templates for hiring, onboarding, mentoring, and planning built for Data and Tech Leads. As a paid subscriber, you get 50% off with code PAIDSUBS-03PL.***



50+ Data Science Templates

JOSE PARREÑO GARCIA · APRIL 8, 2025

[Read full story](#)

And as always, thank you for supporting the newsletter — your support is what keeps it independent, practical, and focused on real-world impact.

OpenAI and Anthropic are becoming consultancies (and you shouldn't be worried about it)

Labs are building deployment arms. Here is who they are targeting, why the consulting layer always returns, and what Data leaders should actually do.

JOSE PARREÑO GARCIA

JUN 18 • PAID • DRAFT



READ IN APP ↗

There has been so much happening this first half of 2026 with OpenAI, Anthropic, Cursos and Chinese model competitors releasing multiple frontier models, agentic harnesses and pricing changes, that there has been a set of news which have been lost under the radar. Around March and April, both OpenAI and Anthropic signed acquisition and set up 2 new companies, which effectively, will mean that both OpenAI and Anthropic will have a consultancy firm under their portfolio.

Read that again.

Both frontier models are betting (with a lot of money) that the future value of AI is on helping others deploy the same AI they build. They are building an AI deployment consulting layer (embedded engineers, partner networks, PE-backed services arms) because they have run into the same wall every major enterprise technology wave runs into.

The model works, but most organisations do not yet know how to absorb it. Someone has to close that gap. And right now, the labs have decided to be part of who does (whilst probably making a future business model out of this).

However, the few headline articles that covered the creation of these consultancy firms scratched just the surface. There is more than to it than

OpenAI's and Anthropic's acquisitions of consultancy firms, therefore, understanding who they are targeting, what the economics look like, and what it signals about the industry — is more useful than the consultancy headline.

This post works through it.

What will we cover in this post?

- **What did OpenAI and Anthropic actually launch?** The concrete structures behind the announcements, in plain terms — not the journalism shorthand.
- **What is a forward-deployed engineer in AI deployment?** The role defined, the Palantir lineage traced, and why the last mile is longer than it looks.
- **How are OpenAI and Anthropic approaching this differently?** OpenAI's vertical ambition versus Anthropic's ecosystem bet — 2 strategies, one structural problem.
- **Who are OpenAI and Anthropic targeting, and who are they not?** The actual customer profile, and why your internal data team is probably not in scope.
- **Why does the consulting layer always return?** Cloud, ERP, RPA, Salesforce — the pattern that repeats every time enterprise technology gets powerful enough to matter.

- **Is this a scalable product-learning loop or just another expensive human layer?** What the Salesforce, ServiceNow, and Palantir numbers say about the economics.
- **What Data and Tech leaders should do about AI deployment consulting.** 3 concrete things to do with this information.

Let's get started!

Senior Data Science Lead is a reader-supported publication. To receive new posts and support my work, consider becoming a free or paid subscriber.

✓ Subscribed

What did OpenAI and Anthropic actually launch?

Both labs launched formal AI deployment consulting arms in May 2026.

OpenAI created a majority-owned [Deployment Company](#) with \$4B+ in investment and ~150 Forward Deployed Engineers via the Tomoro acquisition (an applied AI consulting and engineering firm). OpenAI described the company's purpose as helping organisations "build and deploy AI systems" across important work, with FDEs (forward deployed engineers) leading complex end-to-end production deployments with strategic customers.

May 11, 2026 Company

OpenAI launches the OpenAI Deployment Company to help businesses build around intelligence

OpenAI has agreed to acquire Tomoro, giving the OpenAI Deployment Company experienced Forward Deployed Engineers from day one.

Screenshot from OpenAI's official press release

Alongside this, OpenAI launched [Frontier Alliances](#), a formal partnership with [BCG](#), [McKinsey](#), [Accenture](#), and [Capgemini](#), to define strategy, integrate systems, redesign workflows, and scale deployments globally. In their press-released, OpenAI coupled this alliance with top consulting firms with the creation of DeployCo: *“Our partners will work alongside OpenAI’s Forward Deployed Engineering (FDE) team, combining OpenAI’s research and product expertise with deep transformation experience and global delivery teams. Each partner is investing in dedicated practice groups and building teams that will be certified on OpenAI technology. OpenAI will support them with technical resources, roadmap insight, and access to our product and research teams.”*

On the other hand, [Anthropic announced a separate enterprise AI services company](#) backed by Blackstone, Hellman & Friedman, and Goldman Sachs, with additional investors including Apollo, General Atlantic, GIC, Leonard Green, and Sequoia — total commitments of approximately \$1.5 billion.



Screenshot from Anthropic's official release statement

Applied AI engineers from Anthropic will work alongside the new company's engineering teams to identify where Claude can have the most impact, build custom solutions, and support customers over the long term. Anthropic was explicit that this complements rather than replaces its [Claude Partner Network](#), which already includes [Accenture](#) (30,000 professionals trained on Claude, a joint Centre of Excellence) and [Deloitte](#) (Claude available across a 470,000-person global network).

The timings and the goals were similar.

Both labs announced formal enterprise deployment capacity *within weeks* of each other. [TechCrunch reported](#) that both announcements were driven by the same observed problem: enterprises were accumulating AI pilots without converting them to production. The capability was available, but, the deployment was stuck.

The goal is of course is to *“raise money from alternative asset managers to create new channels for enterprise AI deals. The ventures will presumably get*

preferred sales access to their investors' portfolio companies, while the investors will capture more value from any resulting contracts.”

But, don't think about OpenAI and Anthropic consulting as the historical consulting model.

The journalism shorthand headlines of “labs becoming consultancies”, compress too much information into a single sentence. The reality is that these deals resemble more *“embracing the forward-deployed engineer (FDE) model popularized by Palantir”*, compared to the classical consulting services engagements. The work will be engineering-led, workflow-specific, and explicitly designed to seed recurring model consumption.

What does this look like in practice though? Let's cover that in the next section.

What is a forward-deployed engineer in AI deployment?

A forward-deployed engineer (FDE) is an engineer who embeds directly inside an enterprise customer's environment to build, deploy, and iterate on AI systems in context. This is in contrast to SaaS services selling a platform and leaving deployment to the customer.

Palantir popularised the model, sending engineers to live inside organisations, map real workflows, and feed what they learned back into the platform. *Palantir's AIP Bootcamp* is the cleanest public expression of this —

a structured 5-day engagement where customers go from “0 to use case”, developing initial workflows, onboarding users, and preparing for rollout.

OpenAI and Anthropic are now running the same model for enterprise AI deployment, with FDEs owning complex end-to-end production deployments from day one.

Anthropic shares a very realistic example in their announcement of what a Forward Deployed Engineer engagement could look like.

When I first read Anthropic’s announcement, the language that jumped out was not the valuation or invested money (which is a truck load). It was “Applied AI engineers will work alongside the firm’s engineering team.” That language describes an **operating model**, not a consulting pitch. Below is an example that they share.

“A typical engagement starts with a small team working closely with the customer to understand where Claude can have the biggest impact. From there, the company’s engineers—alongside Anthropic Applied AI staff—will develop Claude-powered systems tailored to each organization’s operations.

Consider a multi-site healthcare services group, like a network of physician practices. Clinicians spend hours each day on documentation, medical coding, prior authorizations, and compliance reviews. An engagement might begin with the company’s engineering team sitting down with clinicians and IT staff to build tools that fit into the workflows that staff already use. The clinicians know where time disappears in a shift and what good patient care

actually requires. The company's engineers build around that knowledge, allowing clinicians to devote more time to patient care"

But will the FDE model work for OpenAI and Anthropic?

Well, [MarketWatch's reporting on Palantir](#) is worth reading carefully: the model worked, but it took more than a decade to mature, it was expensive, and imitators risk misunderstanding it. What made Palantir successful was the feedback loop they created: *"The company developed its products by sending pods of four to five engineers to work with customers for months at a time. FDEs took data from messy spreadsheets, legacy systems and handwritten memos to create a digital replica of an enterprise. Vast amounts of scattered information were categorized and mapped. The result was a continuously updated and standardized system of record across an entire organization. This master data layer, which Palantir called Ontology, became the foundation upon which all of its other applications were built."*

If OpenAI and Anthropic are partnering with classical consulting firms (or replicating them), but not applying FDE as initially designed by Palantir, then maybe the engagement might not be as successful as the hype being created. For example, an FDE who merely customises one-off workflows is a glorified contractor. But, an FDE who converts repeated deployment patterns into reusable platform capabilities is a compounding asset. The job title is easy to copy. The operating model is not.


'People are slapping the FDE title on sales engineering without really understanding how Palantir did FDE or why it worked so well.'

— *Hex Technologies CEO Barry McCardel (ex-Palantir too)*

The FDE hype is real

[The Wall Street Journal](#) noted that FDE-style roles are now proliferating across the AI industry — at labs, cloud providers, enterprise software companies, and consultancies. [The reddit community is also seeing this too.](#)

The titles vary wildly: AI Deployment Engineer, Applied AI Architect, Solutions Engineer, Agent Strategist, GenAI Solution Architect...



The underlying role is consistent: deep enough to build in the customer's environment, broad enough to understand the business workflow, and committed enough to the feedback loop to make the deployment matter beyond one engagement.

For OpenAI, the acquisition of Tomoro seeded that headcount from day one rather than building it from scratch. [Business Insider reported](#) that OpenAI's

FDE team was explicitly intended to build scalable product playbooks, not chase service revenue. The intent, at least, is to run the Palantir model correctly (only time will tell if they did this right).

How are OpenAI and Anthropic approaching this differently?

OpenAI is building vertically whilst Anthropic is building horizontally. Let me expand on these topics.

OpenAI wants to own the enterprise engagement end-to-end

If we recap, OpenAI owns DeployCo and established co-delivery Frontier Alliances with BCG, McKinsey, Accenture, and Capgemini. I think OpenAI is quite ambitious. OpenAI wants to own the enterprise engagement end-to-end — or at least the highest-value parts of it — rather than hand the keys to partners and hope for the best. It's the closest to the Palantir model, but with the punch of having direct access to worldwide consultancy firms and their clients.

Anthropic's strategy looks like a portfolio distribution.

Anthropic's services company is backed by PE firms (Blackstone, Apollo, Goldman Sachs) whose portfolio companies become the natural first customer base. Blackstone, Apollo, and Goldman Sachs *own* positions in hundreds of large enterprises. The implication is that Anthropic's deployment arm is not being sold cold into the market — it is being pre-distributed into

portfolio companies that already have a financial relationship with its backers. That is a channel strategy disguised as an investment announcement.

Anthropic is also mass-training partner workforces rather than building its own delivery army. 30,000 Accenture professionals trained on Claude. The full Deloitte global network with Claude access. These are not “API agreements”. They are the closest to making Anthropic the model layer that GSI delivery capacity defaults to. Most importantly, even when Anthropic engineers *are not in the room*.

The *surface* story is that both labs are doing the same thing. But this is not quite right.

The strategic difference maps to a fundamental question each lab is implicitly answering differently. OpenAI is betting that owning the enterprise deployment relationship generates durable platform lock-in that justifies the cost. Anthropic is betting that making Claude the default choice of the partner ecosystem is more scalable than building delivery headcount it would have to maintain.

Neither has proven the economics. [TechCrunch was direct about this](#): these are capital commitments and intent signals. Not proven margins, not repeatable delivery at scale, not demonstrated ROI for the enterprise customers on the receiving end. The announcements are hypotheses, not results.

Who are OpenAI and Anthropic targeting, and who are they not?

Whilst the strategic approach might be different for both frontier labs, the customer profile they are targeting is not that dissimilar. From both announcements, they look like both are targeting large enterprise, regulated industries, and PE-backed portfolio companies.

Who they both target

Large enterprise is the obvious target. These are organisations with large IT budgets, complex workflows, and enough strategic AI ambition to justify a 1M–5M+ (or more) deployment engagement. They also have the procurement processes, security reviews, and legal requirements that make self-serve API access insufficient. An FDE team that navigates those processes is genuinely valuable because the *organisation around* it is hard to change.

Who is Anthropic distinctively targeting

However, there are some differences. Regulated industries are where the economics get interesting. Finance, healthcare, and government deployments carry heavier requirements: auditability, human-in-the-loop controls, data sovereignty, and incident accountability. [Anthropic's Claude for Financial Services](#) and its explicit targeting of mid-sized companies in its new services company both point in this direction. These customers will pay more, move more slowly, and create more durable contracts — which suits the

Palantir-style economics of high-touch deployment converting to long-term platform revenue.

PE-backed portfolio companies are Anthropic's most distinctive target. As explained before, when Blackstone and Apollo back a deployment services company, their portfolio companies become the natural first customer base, where distribution through financial relationships that already exist.

Who are neither of them targeting

What is absent: individual DS teams, mid-market companies without large IT budgets, and internal data functions in organisations that are already running well. Anthropic does carve out mid-sized companies as an explicit target for its new services arm — partly because that is genuinely underserved, and partly as a concession to the GSI partners who own large enterprise relationships and would reasonably object to direct competition.

The important implication for data leaders in other organisations: a \$4B deployment company with 150 FDEs is not coming to redesign your team's internal analytics workflows. The scale, cost structure, and customer profile of these deployments are calibrated for a different tier. The threat to “your data team” is not the FDE model arriving in your building next quarter.



The reason these deployments are calibrated only for large enterprise is the same reason a consulting layer has always formed around complex technology. (See *next section!*)

Why does the consulting layer always return?

The consulting layer returns every time enterprise technology gets powerful enough to matter because the hard work is never the tool or the tech. The really tough part is the organisation around the tool. 6 recurring jobs sit between “the software works” and “the organisation extracts value”:

1. Translation (capability into workflow)
2. Integration (identity, data, systems)
3. Change management
4. Governance
5. Value measurement
6. Political mediation between competing process owners.

(Very human problems right?)

[McKinsey's analysis of ERP transformations](#) is blunt on why these projects are notoriously over budget: connecting powerful software to real organisations requires workflow redesign, data ownership decisions, and change management for people whose jobs are changing. [Deloitte's intelligent automation survey](#) on RPA documented the same failure mode: **ROI did materialise for organisations that moved beyond piloting**, but **most did not**, and the ones that stalled had impressive dashboards full of bot counts (shiny, but useless).

RPA coined the term “automation theatre.” AI is already rehearsing for its own production run.

So, the same way that the internet, the cloud, or machine learning didn't fix the 6 processes above, you shouldn't expect AI to magically do it. In fact, I would argue that, at the very least, it intensifies one of them — governance — because AI agents can act, not just report.

For example, an ERP system stores process logic. But, an AI agent may *execute* parts of it. That distinction changes the accountability model entirely, and it is why regulated industries will demand auditability and human controls that most enterprise pilots have not yet encountered. [UiPath's guidance on scaling automation](#) describes the destination: not more bots, but an operating model with defined roles, governance standards, and internal champions who own the discipline. The AI equivalent is an agent registry, a model risk committee, and prompt version governance. Doesn't really sound very glamorous right?



ConsultingHumor 
@ConsultingHumor

if AI were to replace consultants, clients would have to be good at describing what they need from consultants.

yeah, i think we're good

Is this a scalable product-learning loop or just another expensive human layer?

Let's first start with the hard truth of data of other companies.

Professional services in enterprise software are almost never the profit engine.

They are the adoption infrastructure.

[Salesforce's FY2025 financials](#) show subscription and support accounting for approximately 94% of total revenue. Professional services occupy the remaining 6% and exist explicitly to help customers achieve business results faster — not to generate margin in their own right.

[ServiceNow's FY2025 annual report](#) is even more direct: professional services and other recorded a gross loss percentage of approximately 5%, partly because partner ecosystem spend rose faster than revenue. ServiceNow runs professional services at a loss, knowingly, because the services accelerate platform subscription growth. The margin lives in the subscription. The services are the shovel that digs the foundation.

Palantir is the right analogue for the FDE model specifically. [Palantir's FY2025 report](#) shows revenue of approximately \$4.5 billion and gross margin around 82% — software-like economics built on a decade of forward-deployed engineering. The AIP Bootcamp model — five days from zero to

use case — is Palantir's productised version of what was originally entirely bespoke. The high-touch work became the factory for reusable patterns that now power a scalable platform.

But Palantir took more than a decade to get there. And [MarketWatch's analysis](#) is explicit that the imitators risk getting it wrong: reducing FDEs to glorified sales engineers, running bespoke custom work with no path to productisation, and discovering that the FDE model was expensive to mature even when Palantir did it correctly.

OpenAI and Anthropic are earlier on this curve than Palantir was at the 82% gross margin stage. The hidden cost that matters most is not the salary of the FDEs, it is the **opportunity cost**. Every senior AI engineer embedded inside a Fortune 500 company's workflow is not building the next capability improvement, the next platform feature, or the next reusable deployment pattern. That compounding opportunity cost is what makes the Palantir-or-consulting bifurcation so consequential: if the embedded work generates reusable product learning, the model compounds. If every deployment is custom work with no abstraction layer, the model is just expensive (as most consultancy companies are).

The question of whether this becomes a scalable product-learning loop or another expensive human layer hinges on one thing: whether the deployment teams are genuinely feeding abstraction back into the platform, or whether the enterprise customers' bespoke requirements are consuming all available bandwidth.

That distinction will only be visible in years of financial results.

What Data and Tech leaders should do about AI deployment consulting.

1) The forward-deployed solutions professional is a real emerging career path.

The FDE role is proliferating fast. [The Financial Times reported](#) on the AI deployment staffing wave across labs, cloud providers, and consultancies — and the job descriptions are consistent regardless of the hiring organisation. The core profile: technical enough to build production AI systems in a customer environment, broad enough to understand the business workflow those systems need to fit, and comfortable enough with ambiguity to operate before the playbook exists. This is not a data scientist role. It is not a solutions architect role. It sits between them, with a heavier emphasis on deployment than either.

Data leaders should watch this role for 2 reasons.

1. First, it is a hiring market they may not yet be competing in and losing people to.
2. Second, it is a signal about which skills compound in the current market. The ability to deploy AI reliably in a specific domain context, run evals, own prompt governance, and measure outcomes is more defensible than the ability to fine-tune a model.

A concrete starting point: pull 3 FDE job descriptions from OpenAI, Anthropic, and one cloud provider. Map the skills listed against your current team's

profile. The gaps you find are your hiring and development priorities for the next 12 months — not a wishlist, a gap analysis with real market data behind it.

2) Your data team is not the target, but, the timeline matters.

Labs are not deploying FDE teams into mid-market companies with functional internal DS functions. The economics do not work: a 150-person FDE arm calibrated for large enterprise and PE-backed portfolio companies is not cost-effective at the scale of an internal analytics team.

But, over time, every repeatable workflow that an FDE team builds for a large enterprise today is a candidate for a packaged agent, a vertical template, or a platform feature tomorrow. [McKinsey's State of AI](#) research consistently shows that the gap between pilot and production is where value is won or lost — and the labs are systematically learning how to close that gap at scale. When those lessons become self-serve products, the downstream effect reaches further down the market.

Pick 1 pattern your team deploys today, say, classification, extraction, summarisation, something repetitive and measurable. Ask: in 3 years, if a lab productises this into a packaged template, does your team's value come from running the template or from knowing which template to use and how to evaluate whether it is working?

The answer tells you where to invest right now.

Study the operating model, not just the announcement.

The FDE playbook — embed, build, feed back, codify — is a template for how any Data Science or data function should be thinking about internal AI deployment right now. Not “what model should we buy?” but “who owns evals when the model updates?”, “who owns prompt version governance?”, “what is the escalation path when an agent produces a wrong output that touches a real workflow?”

The labs are showing, in public, the operating model for AI deployment at scale. That model applies inside your organisation too. The teams that build the internal version of this — owning deployment as a discipline, not just a one-time integration — will be the ones that compound their AI investment rather than restart it every time the model changes.

Now, I want to hear from you

The economics and operating model questions here are genuinely open — and I suspect the answers look different depending on where you sit.

- If you are in a large enterprise, are you already engaging with either lab’s deployment arm — or is the GSI partner channel the actual route in? What is driving that choice?
- For those building internal AI deployment capability: are you treating deployment as a discipline with its own ownership, governance, and eval infrastructure — or is it still attached to individual project teams?

- The Palantir-or-consulting bifurcation feels like the pivotal question. Do you think the labs' FDE teams will generate reusable product patterns, or will enterprise customisation demands keep the work bespoke indefinitely?

Share your thoughts, examples, or counterpoints in the comments. I would love to read your opinions 📌

Further reading

If you are interested in more content, here is an article capturing it and organising it all by topics!



All my articles, one convenient place

JOSE PARREÑO GARCIA · OCTOBER 15, 2024

[Read full story](#)

Senior Data Science Lead is a reader-supported publication. To receive new posts and support my work, consider becoming a free or paid subscriber.

✓ Subscribed

References

- [OpenAI — OpenAI launches the Deployment Company](#) — Official announcement of the OpenAI Deployment Company, its investment

structure, partner list, and FDE headcount via Tomoro acquisition.

- [OpenAI — Frontier Alliance Partners](#) — OpenAI's co-delivery partnerships with BCG, McKinsey, Accenture, and Capgemini for enterprise AI deployment.
- [OpenAI — Introducing OpenAI Frontier](#) — The enterprise Frontier platform and its positioning for production AI deployment with strategic customers.
- [Anthropic — Enterprise AI Services Company](#) — Anthropic's announcement of its new services company backed by Blackstone, Hellman & Friedman, and Goldman Sachs.
- [Anthropic — Claude Partner Network](#) — Anthropic's \$100M partner enablement programme and certification structure.
- [Anthropic — Accenture partnership](#) — Joint Centre of Excellence, 30,000 Accenture professionals trained on Claude, industry-specific offerings.
- [Anthropic — Deloitte partnership](#) — Claude made available across Deloitte's 470,000-person global network.
- [Anthropic — Claude for Financial Services](#) — Anthropic's positioning for regulated financial services deployments.
- [TechCrunch — Both labs launching JVs for enterprise AI services](#) — Coverage of the simultaneous announcements and the enterprise deployment gap both are trying to close.
- [TechCrunch — OpenAI calls in the consultants for its enterprise push](#) — Earlier reporting on OpenAI's evolving enterprise strategy.

- [Palantir — AIP Bootcamp](#) — Palantir's structured five-day FDE engagement model, the closest public analogue to OpenAI and Anthropic's deployment approach.
- [MarketWatch — Palantir pioneered the hottest job in tech; copycats may not succeed](#) — Analysis of the FDE model's maturation at Palantir and the risks for imitators.
- [WSJ — AI startups' new old secret weapon: forward deployed engineers](#) — Industry-wide proliferation of FDE-style roles across the AI sector.
- [Business Insider — OpenAI forward deployed engineers accelerate AI projects](#) — Reporting on OpenAI's FDE team structure, headcount targets, and playbook-building mandate.
- [Palantir — FY2025 annual report / gross margin](#) — Palantir revenue (\$4.5B) and ~82% gross margin; evidence that the FDE-to-platform model can generate software economics.
- [ServiceNow — FY2025 financial results](#) — Professional services gross loss of 5%; services as deliberate adoption infrastructure rather than profit centre.
- [Salesforce — FY2025 revenue mix](#) — Subscription and support at ~94% of revenue; professional services as adoption accelerator.
- [BCG — Rethinking B2B software pricing in the era of AI](#) — Customer pressure for outcome-based AI pricing and the revenue recognition complexity this creates.
- [McKinsey — Getting an ERP transformation back on track](#) — ERP transformation failure analysis; the organisational complexity that technology alone cannot solve.

- [AWS — Migration Acceleration Program](#) — Cloud migration programme combining methodology, partner expertise, and financial support; the enterprise technology adoption factory model.
 - [Deloitte — Intelligent automation survey](#) — RPA ROI analysis; benefits materialise for organisations that move beyond piloting with mature delivery capability.
 - [UiPath — RPA Centre of Excellence](#) — Scaling automation requires operating model, governance, and defined ownership — not just more bots.
 - [McKinsey — State of AI](#) — Annual survey on AI adoption; gap between pilot and production as the central value driver.
 - [FT — AI deployment staffing](#) — Reporting on the AI deployment talent wave across labs, cloud providers, and consultancies.
-

Thanks again for supporting the Senior Data Science Lead newsletter.

Your support makes it possible for me to keep writing openly, while also investing time in the more structured, hands-on work — courses, templates, and tools — that help people apply these ideas in practice. You can find them [here](#).



50+ Data Science Templates

JOSE PARREÑO GARCIA · APRIL 8, 2025

[Read full story](#)



Master Data Storytelling

JOSE PARREÑO GARCIA · NOVEMBER 27, 2025

[Read full story](#)

This newsletter exists in its current form because of readers like you. I am genuinely grateful for that.

Thank you.

 LIKE

 COMMENT

 RESTACK

© 2026 Jose Parreño Garcia
548 Market Street PMB 72296, San Francisco, CA 94104
[Unsubscribe](#)

Get the app

 Start writing