

5 AI agents every solopreneur, manager, and executive should hire

A few months ago, I was talking to a friend who runs a small consultancy (who, of course, had spent a significant amount of time with all the ChatGPT, Claude, Gemini and co). I asked her what had actually changed about how she works or runs consultancy projects.

The answer was as disappointing as it was *not* surprising. She told me that her approach to work had improved on the quality she could produce, for example, being able to research her client needs better, brainstorm better solutions, draft better storytelling.

BUT, and here is the big problematic “but”, her throughput hadn’t improved. She is using these new tools as advanced prompting machines, but she is not scaling her work *through* AI systems.

And this is what I am seeing in many people around me. The conversation about AI at work has been dominated by tools — which model, which app, which subscription. But that is the wrong frame.

[QUOTE BLOCK START]

Tools do not change **how** work gets done.

Roles do.

[QUOTE BLOCK END]

This means that most of us, especially those running teams or businesses should really ask ourselves: *"What roles should AI play in my work?"*

This post is about the roles worth hiring. Not the ones with the most impressive titles — those are traps, and I will explain why in a moment. The roles worth hiring are the ones that handle the repeatable blocking work: the research you defer, the draft you procrastinate, the brief that never gets written because you are already in the next meeting.

When you staff them correctly, they do not replace your expertise. They extend what you can do with it. That is the argument I want to make here, and there is plenty of evidence to convince you to take the step of hiring your first AI agents.

What will we cover in this post?

- **What is the AI org chart as an operating system?** A mental model for mapping AI roles to the work that already needs doing across any team or business.
- **What 5 roles can give you a boost?** Intelligence Analyst, Growth Operator, Voice Operator, Builder Partner, and Operating Chief of Staff. What each one does, what inputs it needs, and where human judgment stays irreplaceable.
- **How do you write an AI role's job description before you build anything?** The single most important thing that separates a governed assistant from a vague chatbot habit.
- **Where should you start, by audience type?** The specific first hires that make most sense for solopreneurs, managers, and executives, and one concrete step for each.

Let's get started!

What is the AI org chart as an operating system?


Every small-medium business, every team and every solopreneur have multiple roles to operate, and all these roles are difficult to operate at scale without hiring more people or working longer hours.

The AI org chart is a mindset change where, each of us, can think about these roles as areas where we can actually “hire” new employees to do exactly what we want without our full oversight. And these new “hires”, would be our own AI agents.

But, there is a twist. This operating system does not assume “hiring” AI Agents to replace human expertise. Literally, it's the opposite.

The goal of the AI org chart is to hire agents to do the well-known work that takes a significant amount of time for a human to do. In some instances, the hired agents could even do a better job than a human, but my mental model is that, given enough time and space, the human would do this job as good as an AI agent.

AI agents fill the gaps. They handle the routine, structured, repeatable parts of each function so that the human doing the job can spend more time on the parts that require judgment, taste, relationships, and experience.



(Caption: Image generated with ChatGPT)

The screenshot above is just an example of a possible AI org chart. Not *the* AI org chart that you need to build (we will get there later). However, before we get into details, let me share with you my personal experience building AI agents to bring some light into how useful the AI org chart can be for you.

My own personal example with 3 agents

For context, I lead 3 Data Science teams (around 25 people and 12 projects) and I am also responsible for the people side of things (calibrations, promotions, etc) for the wider Data Science discipline at Skyscanner (around 60 people).

The “before-AI-agents” Jose would have spent a *significant* amount of time reading Slack messages, Confluence docs and researching technical blogs, just to be up to date on what was happening. My work couldn't simply scale further because the day had 24 hours, and I have the habit of wanting to sleep 8 hours, have lunch & dinner and spend time with my family after work.

But, the “after-AI-agents” Jose has a much richer and deeper signal on everything that is happening. I have hired 3 AI assistants to help me:

[EMBED THE FOLLOWING LINK IN SUBSTACK AND IT WILL RENDER A DATAWRAPPER TABLE]

<https://datawrapper.dwcdn.net/CJwww/1/>

These 3 agents have enhanced the way I work to an extent I couldn't have done. I could probably have more agents, but these are the 3 that have genuinely helped handle routine and structured work.

As a counter example, I do NOT have an agent to run the people side of things. I fully retain that workflow. Nor do I have agents for marketing operations, legal insights or other areas that are not part of my day to day work.

Therefore, what I wanted to show is that the AI org chart should be bespoke to each and one of you. Really think about the areas taking a significant amount of your time which you believe are repeatable and well structured. Those are the areas to build your AI org chart around.

What 5 roles can give you a boost?

Now, having said that *you* need to define your AI org chart, I have seen enough patterns emerge around me to understand that there are potentially 5 roles that would provide you with the most value in a short amount of time. So, in this instance, I'll be opinionated about proposing you to hire 5 AI roles.

The 5 roles are:

- **The intelligence analyst** — intelligence and research function
- **The growth operator** — sales, growth, and customer success function
- **The voice operator** — communications and content function
- **The builder partner** — product and engineering function
- **The operating chief of staff** — operations and coordination function

Let's go through each one of them in detail.

What does the Intelligence Analyst do?

The Intelligence Analyst is your research function. It handles market research, competitor tracking, and customer insight synthesis on a regular cadence. It takes

structured inputs (competitor URLs, customer transcripts, support tickets) and produces a recurring brief. The agent surfaces evidence; you interpret what it means.

In most small teams, this work happens in one of two ways. Either someone does an intensive research sprint before a big decision, or it never happens at all. Neither is good. The Intelligence Analyst turns this into a rhythm (say a weekly memo) or something that can be generated at the tip of your keyboard.

Personally, I think it's the right first AI role to hire because the output is concrete, the risk is low, and the value is immediate.

What it does well:

- Tracks competitor websites, pricing pages, and announcement feeds
- Summarises customer interview recordings or transcripts
- Clusters support tickets by theme and extracts recurring language
- Maps market categories from multiple sources
- Answers specific questions against a defined set of inputs

What it needs to work:

- A product or service description it can reference
- A list of competitors with their URLs
- Customer interview transcripts or support ticket exports
- Clear questions you want answered, not open-ended "research everything"

Where human judgment stays essential:

The analyst finds the answers, but you frame the questions and decide what the answers mean. Market strategy — which segment to target, which competitor to take seriously, which trend to act on — remains a human call.

Are there already tools out there that can help me?

If you are a n8n user, 2 templates can get you started. Not that I would use them directly, but check them out (or ask Claude to do it for you) to get inspiration.

- [n8n's AI-powered market research workflow](#) helps you get a 360-degree view of a product idea or research topic by analysing.
- [competitor analysis automation](#) is a multi-AI-agent system that automates the entire process of competitor and market research, providing you with a complete strategic analysis in minutes.

If you prefer simply copy pasting prompts, try [28 different prompts built by Juma](#).

1. Identifying Competitor Strengths

💡 Prompt:

"Analyze the website, product descriptions, and customer reviews of [Competitor X]. Identify their key strengths in terms of product offering, messaging, and brand positioning. Summarize the top three differentiators they emphasize in their marketing."

"Analyze the website, product descriptions, and customer reviews"

The important thing is to try these systems and tailor them to your needs. Once this is done, I would recommend saving a copy of these prompts into a file (say .txt or .md format) and uploading them to Google Drive. Then its very easy to connect your GPT, Claude or Gemini and simply ask it to run the market research on the instructions you have provided and saved.

What does the Growth Operator do?

The Growth Operator covers 2 jobs that share a common purpose: bringing in customers and keeping them.

The first is sales research and qualification. The second is support triage.

They are distinct workflows, but they belong in the same department because the skills and inputs overlap. In both cases, the agent is trying to answer the same basic question: *what is happening, how important is it, and what should happen next?*

The Growth Operator is not there to run growth on autopilot. The better pattern is analyse-first, draft-second, human-approved-before-send.

For sales, that means researching inbound leads, scoring them against your ideal customer profile, producing a short account brief, and drafting a personalised first message. For support, it means classifying incoming requests, drafting responses for routine queries, and escalating anything sensitive.

What it does well:

- Enriches inbound leads with company, role, market and account context
- Scores leads against your ideal customer profile
- Produces short account briefs before a human reviews the opportunity

- Drafts personalised first messages based on known context
- Flags missing information or low-confidence recommendations
- Classifies support tickets by urgency, topic and risk
- Drafts responses for routine support questions
- Escalates billing, legal, refund, angry customer or reputational-risk cases

What it needs to work:

- A clear ideal customer profile
- A lead list, form submission flow or inbound source
- Examples of good outreach written in your voice
- Clear qualification criteria for what makes a lead worth pursuing
- A support knowledge base or FAQ library
- A taxonomy for support ticket types and urgency levels
- Explicit escalation rules for sensitive cases
- A human approval step before external messages are sent

Where human judgement stays essential:

The Growth Operator can research, classify, score and draft. The human still decides whether a lead is genuinely worth pursuing, whether the timing is right, and whether the message feels appropriate. Fit, tone and prioritisation remain human calls.

The same applies to support. The agent can tell you that a customer is angry, that the issue involves billing, or that the topic has appeared several times this week. But whether the right response is an apology, a refund, a call, a product escalation or simply a clearer explanation should remain with the person who understands the relationship.

There is also a scaling tension worth naming. The research-first, draft-second gate works beautifully for twenty leads a week. At two hundred, the bottleneck can move from the agent to the human reviewer. That is not a reason to remove the gate. It is a signal to tighten your qualification criteria and reduce the volume that reaches the review queue in the first place.

The gate is not the problem. Unclear qualification criteria are.

Are there already tools out there that can help me?

Yes. Clay and Relevance AI are probably the best-known example for sales research and enrichment. They help teams enrich lead lists, research accounts, identify buying signals and build personalised outbound workflows. The useful pattern is not “let the AI spam everyone”. The useful pattern is: give the system a lead, let it gather context,

produce a useful brief, draft a message, and keep a human in the loop before anything gets sent.

I am not advocating buying their services, but I am highlighting that there are companies out there using AI for this. You might not need the full suite, but I am pretty sure you can get inspired by their flow and build relevant prompts for your agents to use. A few [simple ones that you could try can be found in tomba.io](https://tomba.io). If you want a more detailed one, I proposed one below:

[CODE BLOCK]

You are a Growth Research and Qualification Agent.

Your job is to help qualify inbound leads and prepare human-approved outreach. You do not claim certainty where evidence is weak.

You produce a structured brief for human review.

Inputs

Company: {{company_name}}

Website: {{company_website}}

Contact: {{contact_name}}

Contact role: {{contact_role}}

Lead source: {{lead_source}}

Known context: {{known_context}}

Our product/service: {{product_description}}

Ideal Customer Profile: {{icp_definition}}

Qualification criteria: {{qualification_criteria}}

Examples of good outreach in our voice:

{{outreach_examples}}

Task

Research and assess this lead using only the provided information and any connected research tools available.

Produce the following output:

1. Account summary

A short summary of what the company does, who they serve, and why they may or may not be relevant.

2. ICP fit score

Score from 1 to 10.

Include:

- Score
- Confidence level: High / Medium / Low
- Main reasons for the score
- Evidence used
- Missing information

3. Buying signal assessment

Identify any visible trigger events, pains, growth signals, hiring signals, technology signals, or market signals.

Classify each as:

- Strong signal
- Weak signal
- No clear signal

4. Suggested next action

Choose one:

- Prioritise now
- Add to nurture
- Research further
- Disqualify

Explain why.

5. Personalised outreach draft

Write a short first-touch message.

Constraints:

- Professional and conversational
- No hype
- No false familiarity
- No unsupported claims
- No more than 120 words
- Personalised using only verified context
- Clear, low-pressure call to action

6. Human review checklist

List what a human should verify before sending.

7. Final recommendation

Give a concise recommendation in one sentence.

Remember: your role is research-first, draft-second, human-approved-before-send.

What does the Voice Operator do?

The Voice Operator is a governed content function. It takes rough notes, voice memos, or briefs and turns them into first drafts against a written voice guide. The key inputs are your best previous writing as calibration and a voice guide that specifies your register, banned phrases, and what good looks like. Without those inputs, the output sounds like every other AI-assisted content on the internet. With them, it sounds like you.

Most solopreneurs already use AI for writing. The difference between that habit and a Voice Operator is governance. The Voice Operator is a formal content function — inputs, voice guide, style constraints, output format, and a clear human review step before anything publishes.

What it does:

- Turns rough notes or voice memos into first drafts
- Repurposes long-form content for different channels
- Enforces style and tone against a written voice guide
- Produces newsletter outlines, internal updates, investor memos
- Offers alternative hooks, headlines, or framings before you commit

What it needs to work:

- Your best previous writing as calibration material
- A voice guide (even a short one: register, banned phrases, examples of good and bad)
- The intended audience and channel for each output
- A clear output format so it does not guess

The inputs are the difference between an agent that sounds like you and an agent that sounds like every other AI-assisted content on the internet. Without a voice guide, the agent defaults to an average of everything it has seen. That average is recognisable and forgettable.

Where human judgment stays essential:

The voice is yours. The claims are yours. The final editorial call is yours.

The operator can produce a draft that sounds like you. It cannot decide what you actually believe, which stories are worth telling, or whether this piece is ready to publish. Those decisions stay with you.

Are there already tools out there that can help me?

There are plenty of newsletters and blogs with prompts that help you get started with coming up with a voice guide. [I have myself a voice_critic.md for when I write my newsletter](#). It helps me stay grounded and not deviate from my style too far. As you can see, it scores my articles from 1 to 5 against 2 dimensions: voice score and audience score.

What does the Builder Partner do?

Coding agents are the most mature AI role in production today. The tooling has reached a real threshold, as implementation workflows that would have required a junior developer two years ago are now accessible to anyone who can write a clear spec.

But the Builder Partner is not just for technical people. There are two distinct paths. [Claude Code](#), [OpenAI's Codex cloud agent](#), [GitHub Copilot](#), and [Cursor](#) all support both — whether you are writing product specs or implementing code.

The non-technical path: product spec writing

Turn a rough idea into a PRD. Define user stories and acceptance criteria. Break a product concept into implementation tasks. Identify the risks and non-goals before a line of code is written.

This is where managers and executives who are not developers can get real value from the Builder Partner — not by having it write code, but by having it structure thinking before anyone touches a codebase. A well-written spec brief saves hours of back-and-forth with developers.

The technical path: implementation partner

For technical builders, the Builder Partner can implement bounded tasks, write and run tests, review diffs, flag security or edge-case risks, and open a PR for human review. It is not a senior engineer. It is a capable implementer for well-scoped, clearly specified tasks.

The operating rule is simple: it produces work, humans review and approve it.

What it needs:

A repo with context (a `CLAUDE.md` or `AGENTS.md` file that explains conventions, constraints, and what good code looks like here). Examples of good PRs. Acceptance criteria written before implementation starts. Tests that define what "correct" means.

Where human judgment stays essential:

Architecture. Design decisions. What ships and what does not. The builder owns the implementation; the human owns the design and the consequences of the choices.

I really believe this is a game changer for the non-technical folks. In the past, there would have been 40 backs-and-forths on defining what a product manager wanted or the fix that a customer was proposing.

Today, the paradigm changes big time. A non-technical person can describe in English what they want and produce professional product specs. Then, not only can they send the specs to the developers, but they can also build a mini prototype to walkthrough the details with the developers!

What does the Operating Chief of Staff do?

The Operating Chief of Staff is the most overhyped and underbuilt role in the current AI conversation.

Overhyped: because every demo of a "fully autonomous AI chief of staff" glosses over the fact that real chiefs of staff manage stakeholder politics, institutional memory, and judgment calls that are impossible to encode in a system prompt.

Underbuilt: because the narrower, more down-to-earth version is genuinely useful and much easier to build than the marketing suggests. I am talking about meeting prep, inbox summary, project tracking, follow-up drafts, etc.

Start with the morning brief. Every morning:

1. Review today's calendar and flag meetings that need prep
2. Pull prior notes and decisions relevant to those meetings
3. Identify follow-ups due today from the previous week's commitments
4. Draft a suggested priority order for the day
5. Flag anything that has changed since yesterday

This is not glamorous but, given my personal experience, it is also the work that most managers and executives do poorly or skip entirely because it requires focused attention at the start of the day, which is exactly when the inbox has already taken over.

The difference between this role working and not working is almost always the same thing: escalation rules and real system access. Without them, the morning brief surfaces noise. The agent cannot know what changed unless it is connected to the systems where things change. With real calendar access, prior meeting notes, and a task tracker, the brief becomes genuinely useful. The design decision that separates the two is simple: connect the agent to the live systems, or do not build the role yet.

What it needs:

Calendar access, notes from prior meetings, a task or project tracker, and explicit escalation rules for what requires human decision-making before the agent acts.

And it also needs to be built gradually. Meeting prep first. Inbox triage second. Follow-up drafts third. Autonomous email drafting — if ever — much later, after you have calibrated its judgment against your own for months.

Where human judgment stays essential:

The human still runs the meeting. Still manages the relationship. Still decides what is worth escalating and what can wait. The politics, the history, the understanding of why a stakeholder behaves the way they do — none of that lives in a calendar invite. Let the agent prepare context; you still navigate the people.

Are there already tools out there that can help me?

If you have 1 hour, I would watch [Lenny's conversation with Webflow's CPO](#) about building an AI chief of staff. The case is instructive because it starts narrow — email triage and meeting prep — and earns its way into more complex territory over time.

Hopefully, you have been inspired by these 5 roles and seen how the community is already building them using AI agents. I also understand that this can be daunting if you haven't done this before, so I'd like to introduce a couple of ideas to get you started.

How do you write an AI role's job description before you build anything?

If you cannot describe exactly what you want from an AI agent, then you are not ready to build one. Treat it as if you were hiring for a junior or senior member to add to your team. Or, an even better analogy is that you would not hire a contractor and say "help out with the business." You would tell them exactly what you need, what good looks like, what is out of scope, and who to escalate to when something goes wrong. An AI role deserves the same clarity.

Therefore, before building any of the five roles above, write a job description. It does not need to be long. It needs to cover:

Field	What to include
Goal	One sentence: what this role exists to produce
Inputs	What data, documents, or context the role receives every time it runs
Allowed sources	Which URLs, databases, or files it can draw from
Allowed tools	What actions it can take (search, read, write, send — be explicit)
Forbidden actions	What it must never do without human approval
Output format	Exactly what it produces — a doc, a table, a draft, a summary
Escalation triggers	When it must stop and surface to a human rather than continue
Human owner	One named person who reviews, approves, and owns the output

Quality bar	What does a good output look like? Include one or two examples.
-------------	---

Here is what this looks like in practice, using the Intelligence Analyst as an example:

- **Goal:** Produce a weekly market brief covering competitor moves, customer language, and emerging themes relevant to [product/service]
- **Inputs:** Competitor website URLs, recent customer interview transcripts, support ticket summaries from the past 7 days, list of 3–5 questions from last week's brief
- **Allowed sources:** Listed URLs only, no social media speculation
- **Allowed tools:** Web search, summarisation, document writing — no sending or publishing
- **Forbidden actions:** Do not state conclusions about market strategy. Do not assert trends without citing a source. Do not send output to anyone.
- **Output format:** A structured brief with four sections: What changed, What customers said, What assumptions to revisit, What sources support this
- **Escalation triggers:** If a major competitor announcement appears that was not in the prior brief, flag it prominently before summarising
- **Human owner:** Jose — reviews every Monday, decides what to act on
- **Quality bar:** A good brief cites its sources, distinguishes strong signals from weak ones, and includes at least one thing that surprised me. A bad brief asserts things without evidence or restates what I already knew.

The job description does not go into a document and get forgotten. It becomes the system prompt or the workflow configuration for that role. Every time it runs, it runs from that brief.

Write this before you build anything.

Where should you start, by audience type?

Where you start depends on what you do. My suggestion is the following:

Solopreneurs

Start with the Voice Operator and the Intelligence Analyst.

- The Voice Operator because content and communication are typically where solopreneurs have the most work pending and the least time. A governed Voice Operator — one with a real voice guide and clear inputs — compounds quickly. Each piece of content becomes a reference for the next.
- The Intelligence Analyst because research is usually the first thing cut when time is short. A weekly market brief costs you thirty minutes to set up and gives you thirty minutes back every week, indefinitely.

Managers

Start with the Operating Chief of Staff and the Growth Operator.

- The Operating Chief of Staff because the friction in most management roles is not the decisions — it is the preparation required to make them well. Meeting prep, follow-up tracking, and project monitoring are all high-frequency, structured tasks that AI handles reliably.
- The Growth Operator because managers are often the first to feel the gap between the pipeline they have and the pipeline they need. A sales research function that costs no headcount and runs on existing lead data is a fast win.

Executives

Start with the Intelligence Analyst and the Builder Partner.

- The Intelligence Analyst because strategic decisions made without current market intelligence are just expensive guesses. A governed research function that produces weekly briefs transforms the quality of planning sessions without adding a headcount line to the budget.
- The Builder Partner because the cost of slow product development compounds. Even a non-technical executive benefits from having product specs that are more structured, risks that are surfaced earlier, and developers who spend less time on ambiguous requirements.

Closing thoughts: the art of the possible

The org chart framing is a practical design pattern. You are building roles — governed, scoped, owned. You are not adding features to your workflow. You are **changing** what your workflow can do.

When all 5 departments are running, what changes is the ceiling on what your working hours can produce. A solopreneur operates with the research depth of a small team. A manager walks into meetings with context instead of catching up in them. An executive makes strategy decisions with current intelligence rather than last quarter's memory.

None of this replaces the judgment you have built. It makes that judgment more valuable.

- The Intelligence Analyst brings more evidence => the evidence improves your calls.
- The Growth Operator qualifies leads faster => the faster qualification means more of your time goes to the customers who actually benefit from your expertise.

That is the art of the possible. Not a shorter working day. A wider range for the same working day.

You still run the operation. But now it runs better.

Now, I want to hear from you

I am curious about 2 things.

1. First, which of these 5 roles feels most immediately applicable to how you work right now? And which one feels furthest from something you would actually build?
2. Second — and this is the question I think matters most — when you imagine these roles running well in your context, what is the work you would spend that recovered time on? What is the thing that keeps getting pushed back because the research is not ready, or the brief is not written, or the inbox digest never happened?

Share your thoughts, examples, or counterpoints in the comments. I would love to read your opinions 🙌

References

- [n8n AI market research workflow](#) — Automated market research pipeline using Groq, OpenAI, and Gmail
- [n8n AI competitor analysis workflow](#) — Competitor analysis reports using GPT-4, Apify, and Google Docs
- [28 market research prompts by Juma](#) — You can copy paste and adapt.
- [Clay lead enrichment](#) — Lead enrichment and GTM research platform used widely in SDR workflows
- [Tomba.ai research prompts](#) - With a comparison across Gemini, Claude and ChatGPT.
- [Claude Code](#) — Anthropic's agentic coding CLI for implementation, review, and testing workflows
- [OpenAI Codex cloud agent](#) — OpenAI's cloud-based coding agent for PR-scoped implementation tasks
- [GitHub Copilot](#) — GitHub's AI pair programming tool for code completion and review
- [Cursor](#) — AI-native code editor with agent capabilities for implementation tasks
- [Lenny's Newsletter — Webflow CPO on AI chief of staff](#) — Interview with Webflow's CPO on building and using an AI chief of staff in practice
- [OpenAI practical guide to building AI agents](#) — OpenAI's recommendations for human oversight and approval gates in agent workflows
- [Anthropic guide to building effective agents](#) — Anthropic's engineering guidance on agent architecture and least-privilege design