Gmail **Jose Parreño Garcia <joparga3@gmail.com>**

[TEST ccb - Paid] Sonnet 5 is only good at one thing (beating Sonnet 4.6)
1 mensaje

Jose Parreño Garcia from Senior Data Science Lead <joseparreogarcia@substack.com>
Para: joparga3@gmail.com
15 de julio de 2026 a las 14:57

View in browser

The Senior Data Science Lead newsletter

Hi! It's Jose from Senior Data Science Lead newsletter, where I write about senior data science leadership, sharing hard-earned lessons from building teams, systems, and real-world machine learning.

Because you are a paid subscriber, you are directly supporting this work — and you get full access to the practical tools built alongside the writing:

- ***The full Data Storytelling course**, focused on turning analysis into insights that land with executives. Because today, anyone can build a chart with AI, but very few people can build the right chart.*

📊 **Master Data Storytelling**
JOSE PARREÑO GARCIA · NOVEMBER 27, 2025
Read full story

- *50+ Notion templates for hiring, onboarding, mentoring, and planning built for Data and Tech Leads. As a paid subscriber, you get **50% off** with code **PAIDSUBS-03PL**.*

💼 **50+ Data Science Templates**

JOSE PARREÑO GARCIA · APRIL 8, 2025

Read full story

And as always, thank you for supporting the newsletter — your support is what keeps it independent, practical, and focused on real-world impact.

Sonnet 5 is only good at one thing (beating Sonnet 4.6)

An article to dissect how much can you trust Anthropic's press release vs the community backlash.

JOSE PARREÑO GARCIA

JUL 15 · PAID · DRAFT

READ IN APP ↗

On the 30th of June, Anthropic released Sonnet 5 and what it should have been a launch to get people excited, seems to have transformed in quite the opposite. When I myself read that *"Sonnet 5 narrows the gap: its performance is close to that of Opus 4.8, but at lower prices"* directly from Anthropic's release blog, I thought that at least Anthropic had started to tackle what I believe is the future of LLMs: smaller and cheaper but capable enough models.

But no, the community did not buy Anthropic's release. The amount of backlash I read in Reddit, Substack and LinkedIn was as if Sonnet 5 was pure garbage. Well I don't buy that narrative. Now, I also totally disagree with Anthropic's sales-pitch putting Sonnet 5 at Opus 4.8 levels.

So, where does this leave us? Is Sonnet 5 a good or a bad model? Or should the question be, is Sonnet 5 a good or a bad model for certain tasks?

That's what I aim to answer in this post. We will cover the actual technical changes implemented in Sonnet 5, dissect Anthropic's launch chart claims and do a fair comparison at the backlash it received.

✔ Subscribed

What will we cover in this post?

- **What is Claude Sonnet 5's agentic upgrade?** What actually changed in the model versus Sonnet 4.6: a new top effort tier, thinking that's now on by default, and the tokenizer swap.
- **What does Anthropic claim about Sonnet 5's benchmark gains?** The launch chart, the official deltas, and what "most agentic Sonnet" actually means in numbers.
- **Is Sonnet 5 actually cheaper at matched intelligence tiers?** The one comparison that survives scrutiny: Sonnet 4.6 medium versus Sonnet 5 medium.
- **Where does Sonnet 5 down?** The online backlash and the behavioural complaints,.
- **Should you trust a launch chart or a viral tweet?** A short closing thought on what to actually do with all of this.

Let's get started!

What is Claude Sonnet 5′s agentic upgrade?

Sonnet 5's agentic upgrade comes down to 3 mechanical changes:

1. A new top effort tier
2. Thinking that turns itself on by default
3. A new tokenizer

What effort levels does Sonnet 5 support that Sonnet 4.6 doesn′t?

Let's start with effort levels, because nothing else in this post makes sense without them.

If you are not familiar about what the effort parameter is, then in a summarised way it is a dial for how long the model is allowed to think and act before it has to commit to an answer (more effort means more reasoning tokens, more tool calls, more self-checking passes, and a slower, more expensive response). Low effort is a model trying to get it right on the first attempt. Max effort is a model given permission to inspect, act, test, and revise as many times as it judges useful. If you are interested in knowing more about thinking levels, check my detailed post about it.

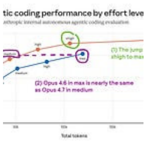

You don't need thinking levels in Claude Code. You need Planning and Goal modes.

JOSE PARREÑO GARCIA · JUL 5

Read full story

Sonnet 4.6 already had this dial — `low`, `medium`, `high`, and `max` — so the concept isn't new. What's new is a level Sonnet 4.6 never had at all: `xhigh`, sitting between high and max, built specifically for long-horizon coding and agentic work that runs for extended stretches. Anthropic's own effort documentation lists `xhigh` as available on Sonnet 5 but not on Sonnet 4.6 (so it's a genuine new tier).

So, a small extra effort level we can play with in Sonnet 5 to tune our workflows.

What is adaptive thinking in Sonnet 5?

I think this is genuine flip. On Sonnet 4.6, a request with no `thinking` field simply ran without thinking, you literally had to opt in. On Sonnet 5, that same request now thinks by default; you have to explicitly pass `thinking: {type: "disabled"}` to turn it off. Manual extended thinking, where you handed the model a fixed `budget_tokens` and it thought exactly that much, is gone entirely on Sonnet 5 — it still worked (deprecated but functional) on Sonnet 4.6, and now returns an error. The model decides how much to think for itself; you no longer get to hand it a number.

Some people might see this loss of control as something not great, but if we do believe that agents will get smarter in the future, then routing thinking efforts should an LLM problem to handle for most tasks.

What does it mean for Sonnet 5 to have a new tokenizer?

Sonnet 5 uses a new tokenizer that maps the same text to more tokens than before — Anthropic's own estimate is a range of roughly 1.0 to 1.35 times, with about 30% as a useful planning figure. Simon Willison ran his own token-counter tests and got sharper numbers: about 1.4x more tokens for English text, 1.33x for Spanish, 1.28x for Python code, and almost no change for Simplified Mandarin.

I used my Claude Token Counter tool to try out the new tokenizer. Here are my results for several larger documents:

Document	Sonnet 4.6	Opus 4.7	Sonnet 5
Universal Declaration of Human Rights (English)	**2,356**	**3,347** 1.42x	**3,341** 1.42x
Universal Declaration of Human Rights (Spanish)	**3,572**	**4,753** 1.33x	**4,747** 1.33x
Universal Declaration of Human Rights (Chinese, Mandarin Simplified)	**3,334**	**3,366** 1.01x	**3,360** 1.01x
sqlite_utils/db.py (4,279 lines of Python)	**44,014**	**56,118** 1.28x	**56,113** 1.27x

So the new token is roughly 1.4x times more expensive for English, 1.33x for Spanish, 1.28x for Python code and effectively the same cost for Simplified Mandarin.

Screenshot taken from Simon Willison blog post

At face value, it seems this is worse comparing Sonnet 4.6 vs Sonnet 5. If we just look at token counts, knowing that the pricing between both versions is the same, then a request just got roughly 40% more expensive.

The tokenizer change also eats into the advertised context window. Sonnet 5's nominal context is still 1 million tokens, same as Sonnet 4.6. But if identical source text now consumes roughly 1.3x as many tokens, then a document that used to fill 769,000 old-token-equivalents of that same 1-million-token window now fills the whole thing. The window is the same size on paper. The amount of your actual document that fits inside it went down. Again, this seems to be something worse.

These are the main changes as per Anthropic's release. I assume more would exist in the internal model training mechanics, but I focused on the official documentation. Now, none of this means much unless we look at the metric benchmarks.

What does Anthropic claim about Sonnet 5's benchmark gains?

Anthropic positions Sonnet 5 as its most agentic Sonnet, built for advanced coding, long-running agents, and professional work.

Do Sonnet 5 scores look strong against Sonnet 4.6?

Yes, they do look genuinely strong. Check the launch chart below.

	Sonnet 5	Sonnet 4.6	Opus 4.8 For reference
Agentic coding SWE-bench Pro	**63.2%**	58.1%	69.2%
Agentic coding Terminal-Bench 2.1	**80.4%**	67.0%	82.7%
Multidisciplinary reasoning Humanity's Last Exam	**43.2%** no tools	34.6% no tools	49.8% no tools
	57.4% with tools	46.8% with tools	57.9% with tools
Computer use OSWorld-Verified	**81.2%**	78.5%	83.4%
Knowledge work GDPval-AA v2	**1618**	1395	1615

On Anthropic's own system-card evaluation:

- Sonnet 5 rises from 58.1% to 63.2% on SWE-bench Pro — a test of fixing real, harder-than-usual GitHub issues across multiple files. It rises from 67.0% to 80.4%
- Terminal-Bench 2.1, a command-line execution benchmark.
- FrontierCode, which grades real pull-request implementation against held-out tests, jumps from 15.1% to 38.8% — more than double.
- AutomationBench, an enterprise workflow-automation test, goes from 5.3% to 13.5%, also more than double in relative terms even though the absolute numbers stay low.

Those 4 are worth highlighting because they share a patternnot really: agentic tasks that reward a model for persisting, checking its own work, and reaching a working end state rather than answering correctly on the first pass. That's consistent with everything the mechanics section just established (more effort, more turns, more self-verification).

Can we trust Sonnet 5 scores benchmark scores against Sonnet 4.6?

⚠️ ***A caveat before trusting these numbers at face value***

Anthropic's own methodology appendix discloses that the Terminal-Bench 2.1 comparison ran Sonnet 5 at `xhigh` effort against Sonnet 4.6 at `high` effort. Which means the comparison is **not the same tier**.

Given that Sonnet 5 `medium` is disclosed as comparable to Sonnet 4.6 `high`, and `xhigh` is a tier Sonnet 4.6 never had access to at all, Anthropic's single biggest headline gain on this chart is built on an effort-level mismatch, not a fair fight between 2 models running at the same setting.

The appendix also doesn't state the effort settings used for Sonnet 4.6 on the other headline benchmarks at all, but it's reasonable to assume the same pattern may hold there too. Likewise, the effort levels for Opus 4.8 isn't mentioned either, but I highly doubt it's on `xhigh` with such small differences compared to Sonnet 5.

What the chart doesn't tell you at a glance is what any of those gains cost per task, or whether the effort level that produced 80.4% on Terminal-Bench is the effort level you'd actually run in production. Those are the 2 questions that determine whether "most agentic Sonnet" translates into "the model you should be paying for." The first place to check is the one comparison built specifically to answer it.

Is Sonnet 5 actually cheaper at matched intelligence tiers?

Sonnet 5 is better than 4.6 by a substantial amount (on Anthropic benchmarks)

We already saw the benchmarks of Sonnet 5 at `xhigh` vs Sonnet 4.6 at `high` above. Let's look at the `high` only numbers though. This will make comparisons fairer.

What you can see is the following:

- Pass rates of Sonnet 5 bump to 80% vs Sonnet 4.6 which are around 75%.
- And they do so at a lower cost per task. Around $7 vs $25.

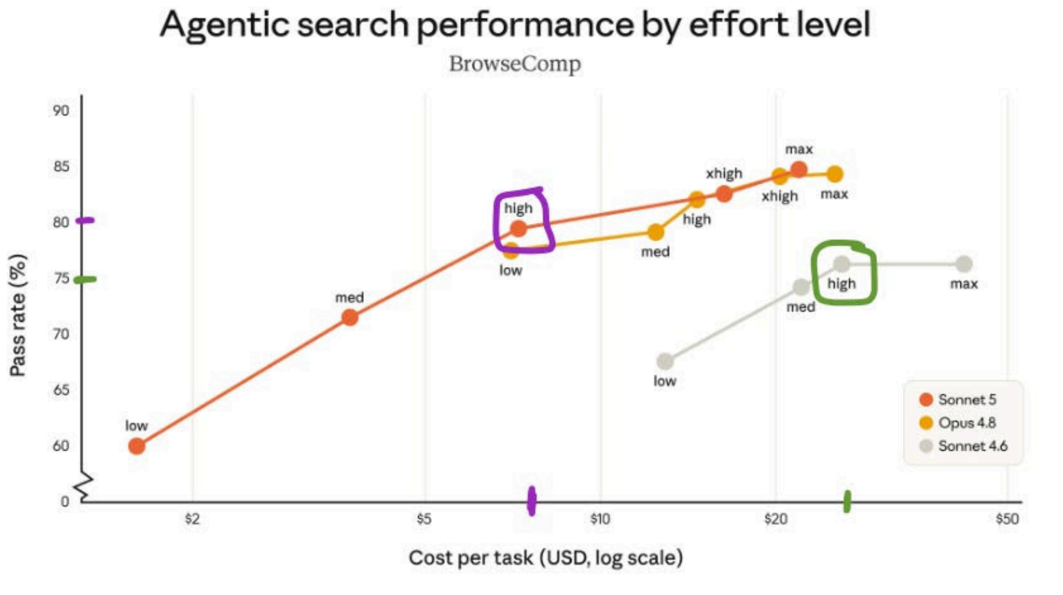

Independent benchmarks contradict that Sonnet 5 is cheaper than Sonnet 4.6

Artificial Analysis found that Sonnet 5 costs 2.29 per task on their Intelligence Index, which is roughly ~2x the price of Sonnet 4.6.

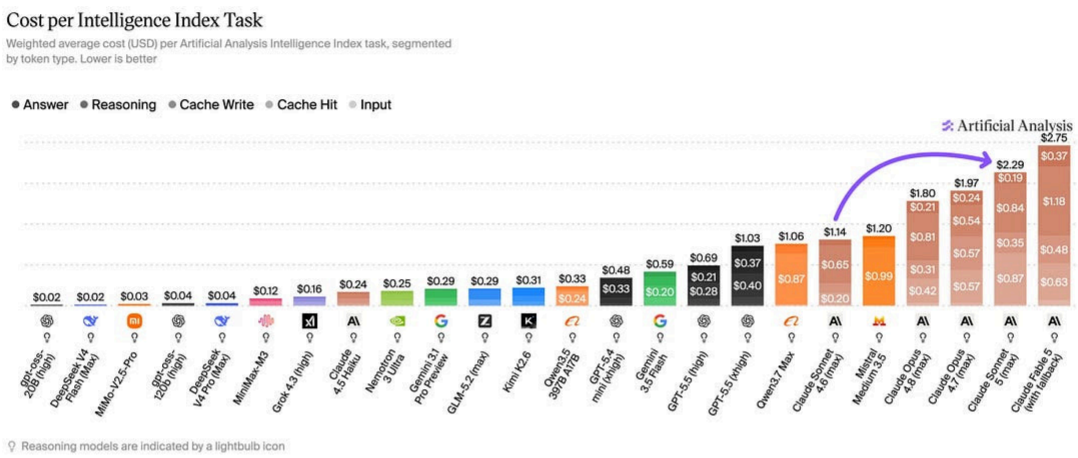

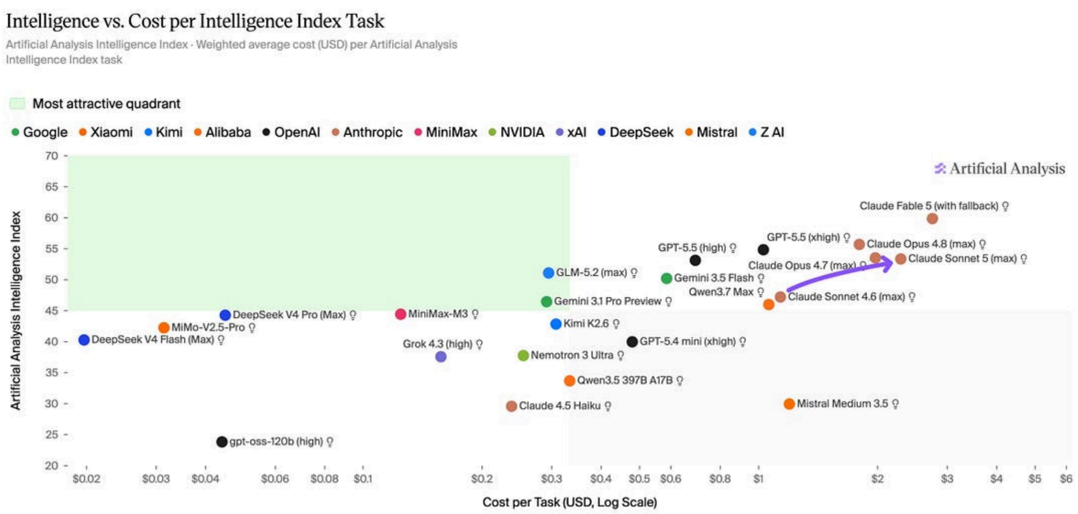

It is true that ArtificialAnalysis.ai is comparing max thinking levels, where Sonnet 5 used roughly 40% more output tokens and around three times as many agentic turns as Sonnet 4.6 based on their benchmarks. This contradicts the cost per task chart from Anthropic, because at max levels, Anthropic claims that Sonnet 5 is cheaper than Sonnet 4.6.

So, who to trust? Probably only your own benchmarks. I guess I am in the middle court. I rarely use max levels and try to run well spec'd tasks with medium. Therefore, I am certain that Sonnet 5 will be better than Sonnet 4.6 given the price difference reported by Anthropic, but I wouldn't take the cost savings per task at face value.

Sonnet 5 is more precise at catching bugs than Sonnet 4.6

CodeRabbit.ai benchmarks shows that Sonnet 5 is better as a coding assistant. It does catch less bugs than 4.6, but it is more precise. It also adds less noise to PR comments.

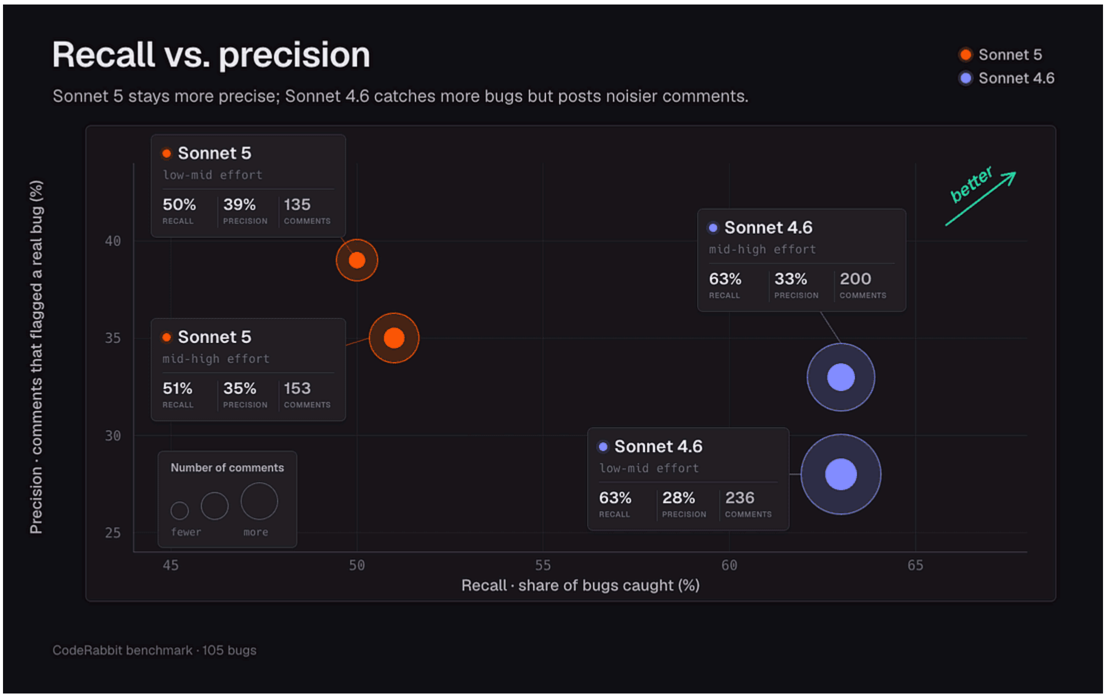

As they mention: *"In plain terms: switch now if you write or ship real software and want a model that tests its own work and sticks with a hard problem until it's solved. Run it at medium effort and you get most of the upside without the top-tier price."*

Sonnet 5 does get close to Opus 4.8

On `max` settings, it is definitely close. Both Anthropic's and Artificialanalysis.ai charts show a really narrow gap in the cost per task and in pass rates (or intelligence indexes).

The only thing I would highlight again is that `max` settings tend to be pretty expensive regardless of the model. So, if we look at other settings, you can see how, despite a cheaper price per million tokens of Sonnet 5, Opus 4.8 matches cost and intelligence at a `low` level. But, at equal `high` settings, Sonnet 5 does a decent job.

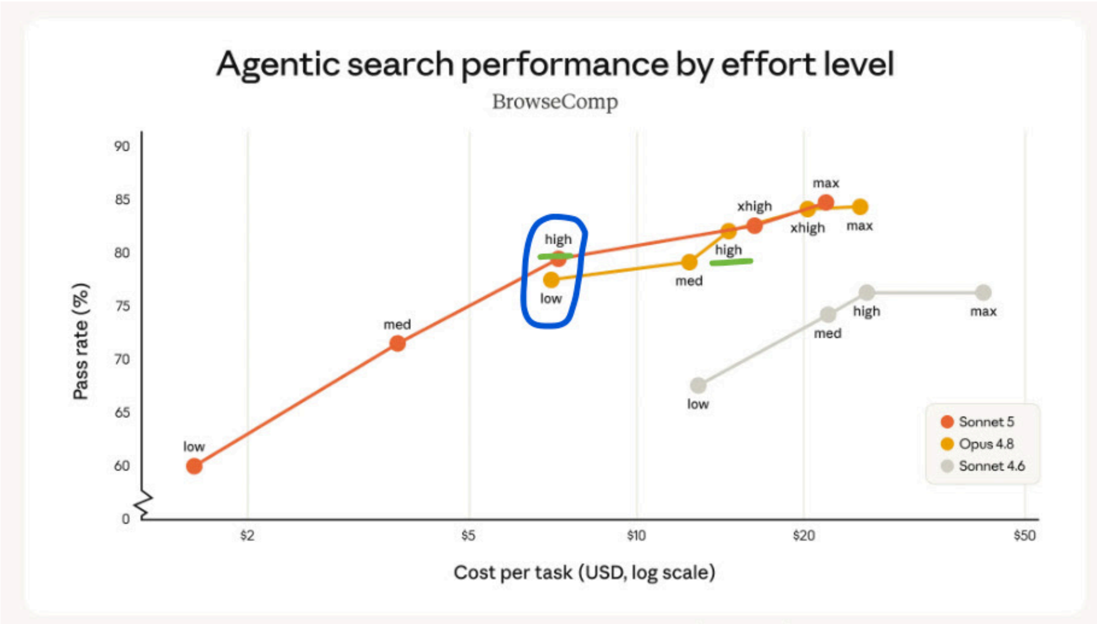

However, the claim must be repeated. It gets *close*. But the cost per task is pretty much the same really. So, don't blindly take the banner news without checking what really matter: the money you are spending for what you are getting back.

Where does Sonnet 5 break?

When I looked at the numbers above, I was surprised by the amount of backlash the model had. But the fact is, that the users are the REAL TEST.

You can measure everything on benchmarks, but benchmarks can be "gamified" or "overfitted". Also, benchmarks dont fully capture the whole user base needs.

And, unfortunately for Anthropic, the community did not like it. I recommend Katie Parrott's blog post on this same topic, and would like to highlight comments from X that she shares.

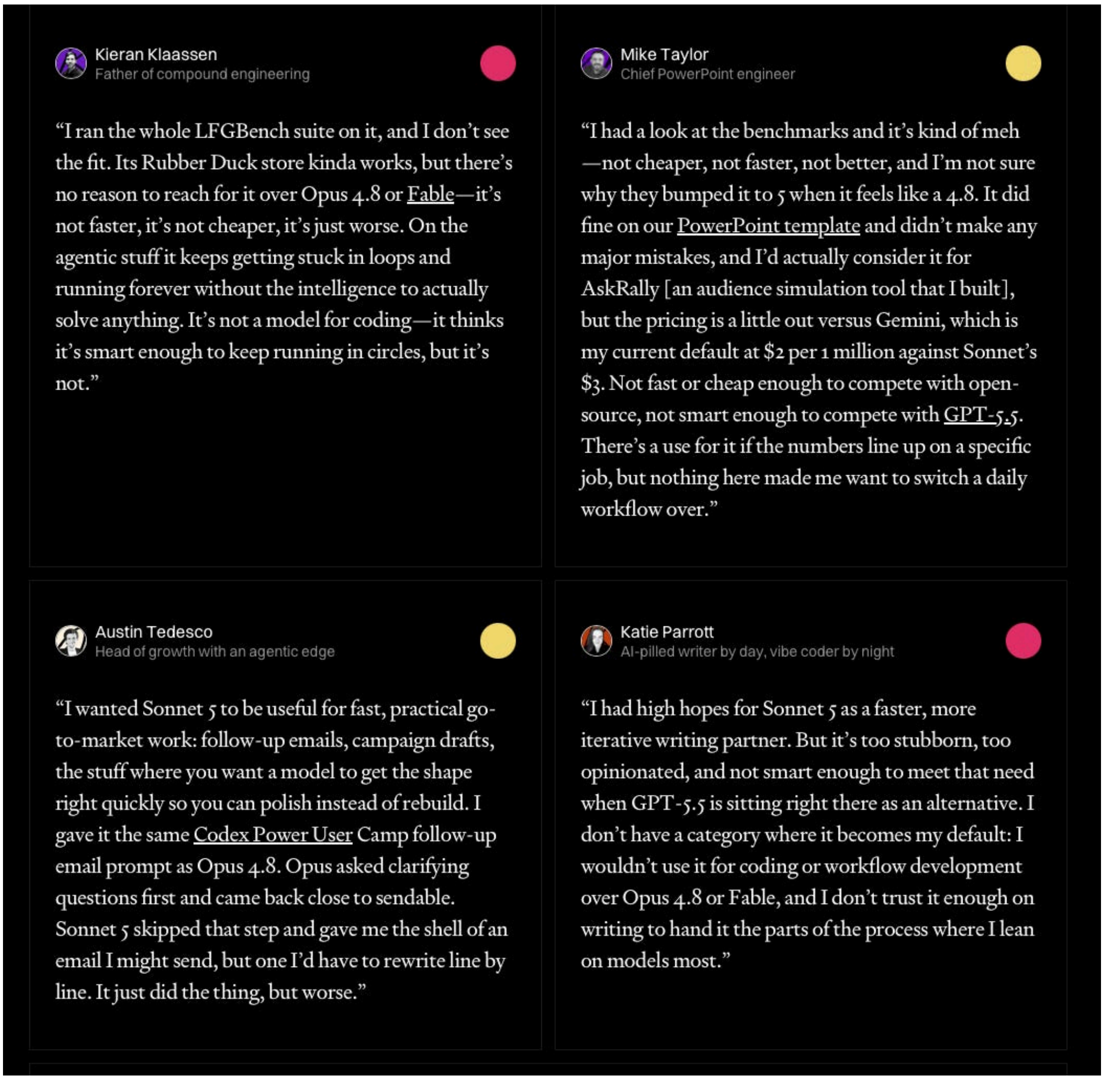

Kieran Klaassen
Father of compound engineering

"I ran the whole LFGBench suite on it, and I don't see the fit. Its Rubber Duck store kinda works, but there's no reason to reach for it over Opus 4.8 or Fable—it's not faster, it's not cheaper, it's just worse. On the agentic stuff it keeps getting stuck in loops and running forever without the intelligence to actually solve anything. It's not a model for coding—it thinks it's smart enough to keep running in circles, but it's not."

Mike Taylor
Chief PowerPoint engineer

"I had a look at the benchmarks and it's kind of meh—not cheaper, not faster, not better, and I'm not sure why they bumped it to 5 when it feels like a 4.8. It did fine on our PowerPoint template and didn't make any major mistakes, and I'd actually consider it for AskRally [an audience simulation tool that I built], but the pricing is a little out versus Gemini, which is my current default at $2 per 1 million against Sonnet's $3. Not fast or cheap enough to compete with open-source, not smart enough to compete with GPT-5.5. There's a use for it if the numbers line up on a specific job, but nothing here made me want to switch a daily workflow over."

Austin Tedesco
Head of growth with an agentic edge

"I wanted Sonnet 5 to be useful for fast, practical go-to-market work: follow-up emails, campaign drafts, the stuff where you want a model to get the shape right quickly so you can polish instead of rebuild. I gave it the same Codex Power User Camp follow-up email prompt as Opus 4.8. Opus asked clarifying questions first and came back close to sendable. Sonnet 5 skipped that step and gave me the shell of an email I might send, but one I'd have to rewrite line by line. It just did the thing, but worse."

Katie Parrott
AI-pilled writer by day, vibe coder by night

"I had high hopes for Sonnet 5 as a faster, more iterative writing partner. But it's too stubborn, too opinionated, and not smart enough to meet that need when GPT-5.5 is sitting right there as an alternative. I don't have a category where it becomes my default: I wouldn't use it for coding or workflow development over Opus 4.8 or Fable, and I don't trust it enough on writing to hand it the parts of the process where I lean on models most."

It is clear that Anthropic probably got very wrong putting Sonnet 5 as a comparison to Opus 4.8. In my opinion, had they simply made the Sonnet 4.6 comparison, the community take would have been more favourable.

One thing is comparing *"hey, for the same accuracy, I get a much cheaper cost per task"*. A very different one is saying that Sonnet 5 max is similar to Opus 4.8 max… because then the community tries it, and it goes badly.

Bleep.co also talks about how users feel that the new security harnesses are hurting Sonnet 5 over previous versions: *"Longtime Sonnet users piled on too. Some said older versions, especially Sonnet 4.6, felt faster, smarter, or less censored on everyday prompts. Others complained about higher refusal rates, slower perceived speed, and what they called AI shrinkflation, the sense that each new release delivers less real improvement than the one before."*

Should you trust a launch chart or a viral tweet?

Neither, on its own.

On one side, a launch chart tells you what a vendor's benchmark suite showed at the effort level that produced the biggest number.

On the other, a viral tweet tells you what one person's workload felt like at whatever effort level they happened to be running, usually without saying which one.

Both are real data points but neither is a substitute for measuring *your* own workload at *your* own effort level.

Run the same test on your own tasks, and make it a real procedure rather than a vibe check. Pick 15-20 representative tasks from your actual workload — not a benchmark suite, *your* workload.

These could be coding tasks or writing tasks. Some can be measured with accuracy numbers, others you can measure with a "liked" or "didnt like". If Sonnet 5 at the matched tier beats your current cost-per-success by a meaningful margin, that's your switch signal. If you have to jump to high effort or above to see a real accuracy gain, price that jump in cost-per-success terms before assuming it's worth it — this post's own evidence says it usually isn't.

The answer might agree with the launch chart. It might agree with the backlash. It will almost certainly disagree with both a little, because your workload isn't Anthropic's benchmark suite and it isn't a stranger's tweet either.

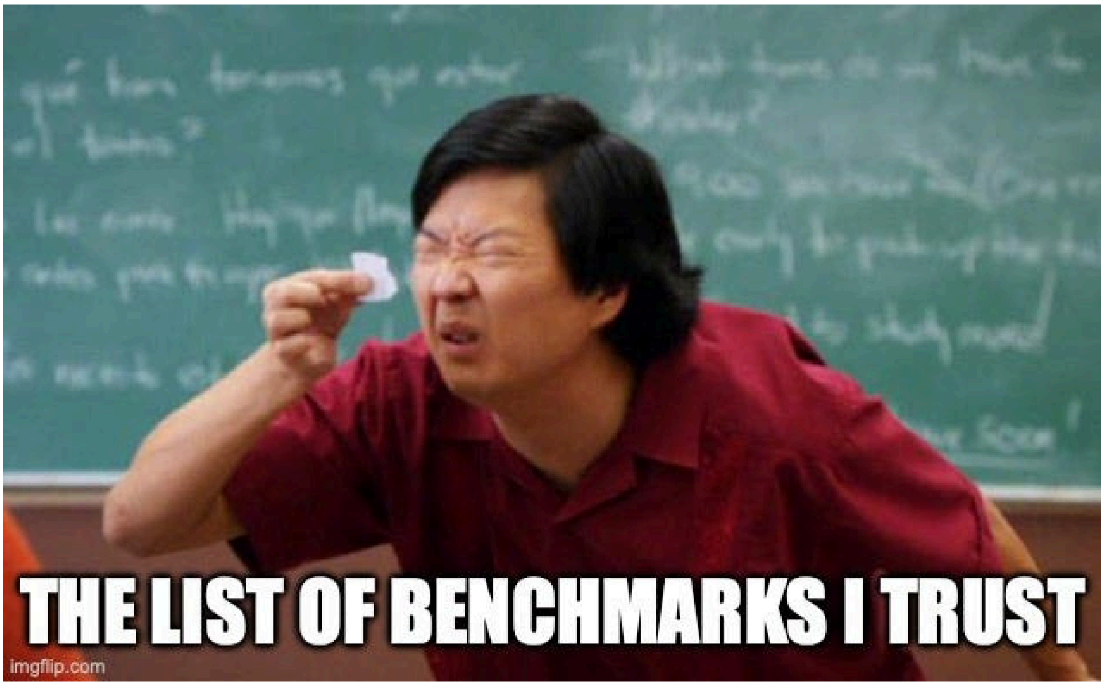

Now, I want to hear from you

I'm genuinely curious about your opinions if you are heavy Claude user and have tried Sonnet 5.

- Which effort level are you actually running in production — and have you checked whether it's the one your cost comparisons assumed?
- Have you measured cost-per-successful-task on your own workload, or only cost-per-attempt?
- Where have you caught a benchmark claim, positive or negative, quietly comparing two different effort levels or model tiers?

Drop a comment or reply — I read everything.

Further reading

If you are interested in more content, here is an article capturing all my written blogs!

All my articles, one convenient place

JOSE PARREÑO GARCIA · OCTOBER 15, 2024

Read full story

References

- Introducing Claude Sonnet 5 (Anthropic) — Anthropic's own launch announcement, source of the "narrows the gap... close to that of Opus 4.8" framing quoted in the opening.
- What's new in Claude Sonnet 5 (Simon Willison's Weblog) — independent verification of the tokenizer change and confirmation that manual sampling controls were removed.
- Effort (Claude API docs) — Anthropic's official documentation on effort levels, source of the confirmed "Sonnet 5 medium is comparable to Sonnet 4.6 at high effort" mapping.
- Claude Sonnet 5: strong agentic performance at a higher cost per task (Artificial Analysis) — the live, independent source for the $2.29-per-task figure and the token/turn consumption behind it.
- Claude Sonnet 5 model page and comparison data (Artificial Analysis) — live pricing and Intelligence Index confirmation for Sonnet 5.
- CursorBench (Cursor) — Cursor's production coding-agent benchmark methodology, confirming 3.1 as the current public version.
- Claude Sonnet 5 Is Not Frontier but Has Its Uses (Don't Worry About the Vase / Zvi Mowshowitz) — a named, non-anecdotal critical take on Sonnet 5's positioning against Opus 4.8 and Fable 5.
- "Ask HN: Is it just me or does Claude / Sonnet 5 sound condescending recently?" (Hacker News) — a single anecdotal thread on tone complaints, cited as a live but low-confidence data point.
- Claude Sonnet 5 review: Should you switch? (CodeRabbit) — independent practitioner data on the construction-versus-review precision/recall tradeoff.
- Bleep.co - Claude Sonnet 5 called a useless flop by fans despite near-Opus benchmarks
- Katie Parrott's blog post - Vibe Check: Sonnet 5—A Model Pitched for Everyone Impresses No One
- What's new in Claude Sonnet 5

Thanks again for supporting the Senior Data Science Lead newsletter.

Your support makes it possible for me to keep writing openly, while also investing time in the more structured, hands-on work — courses, templates, and tools — that help people apply these ideas in practice. You can find them here.

💼 50+ Data Science Templates

JOSE PARREÑO GARCIA · APRIL 8, 2025

Read full story

📊 Master Data Storytelling

JOSE PARREÑO GARCIA · NOVEMBER 27, 2025

Read full story

This newsletter exists in its current form because of readers like you. I am genuinely grateful for that.

Thank you.

LIKE COMMENT RESTACK
