The hardest calls in management: deciding who deserves recognition

5 real stories that challenge how we define, evaluate, and reward impact in tech.

JOSE PARREÑO GARCIA AND RAFA PÁEZ

SEP 27, 2025

3 1 Share

Today's article is an article that Rafa Páez and myself wrote a few months ago and published it in his newsletter The Engineering Leader.

Rafa Páez is one of my favourite authors and I've had great conversations with him. And as a bonus, he is starting to write more and more about AI through his engineering leadership lens. Make sure you subscribe to his newsletter!

The Engineering Leader

The go-to newsletter for high agency engineers and leaders ready to elevate their craft and multiply their impact.
By Rafa Páez

joparga3@gmail.com Subscribe

This article tries to dive into a question that most tech leaders face every performance review: **How do you recognise great performance… when it doesn't always look like output?**

This is a question we have ourselves come back to again and again as leaders.

Not all impact is loud and coupled with shiny numbers. Some impact can be seen immediately, and other has a slow gradual build up. And of course, not all high performers *look* the same.

So how do you compare? How do you evaluate? How do you avoid over-rewarding the most *visible* work and overlooking the most *valuable*?

A few months ago, Rafa Páez and I tried to unpack this in a piece we co-wrote for his newsletter. We looked beyond the surface metrics to ask: **what does holistic performance really look like for engineers?**

If you have ever wrestled with recognising impact across wildly different types of contributions, this one is for you.

Senior Data Science Lead is a reader-supported publication. To receive new posts and support my work, consider becoming a free or paid subscriber.

What will we cover in this post?

Today, we want to share with you 5 real stories where traditional impact measurement falls short. Each story will challenge a common assumption:

1. **The £1M that wasn't fully recognised**. Why a machine learning improvement that generated $1M in revenue wasn't considered strategic.
2. **The £1M that *was* recognised** (**and why context matters**). A parallel example where a similar financial outcome led to much greater recognition.
3. **The problem with recognition-based impact**. How a team that provided company-wide enablement received 10x the visibility but wasn't necessarily creating 10x the impact.
4. **The challenge of measuring slow-moving metrics**. Why long-term improvements often go unrewarded, even when they fundamentally change an organisation's

capabilities.

5. **The amplifier effect: impact that multiplies over time**. How a single contributor reshaped engineering standards and de-risked projects without producing a measurable metric.

Let’s begin!

Story #1: Flights ranking: the £1M that wasn’t fully recognised

> *"A £1M revenue boost sounds like a success story. So why wasn’t it enough?"*

Let’s start with a case that should, on paper, be an easy win.

One of our senior data scientists was working on flights ranking, a core component of how we optimise search results for travellers. The goal was to improve how we rank flight options, leading to better user engagement and, ultimately, higher conversion rates

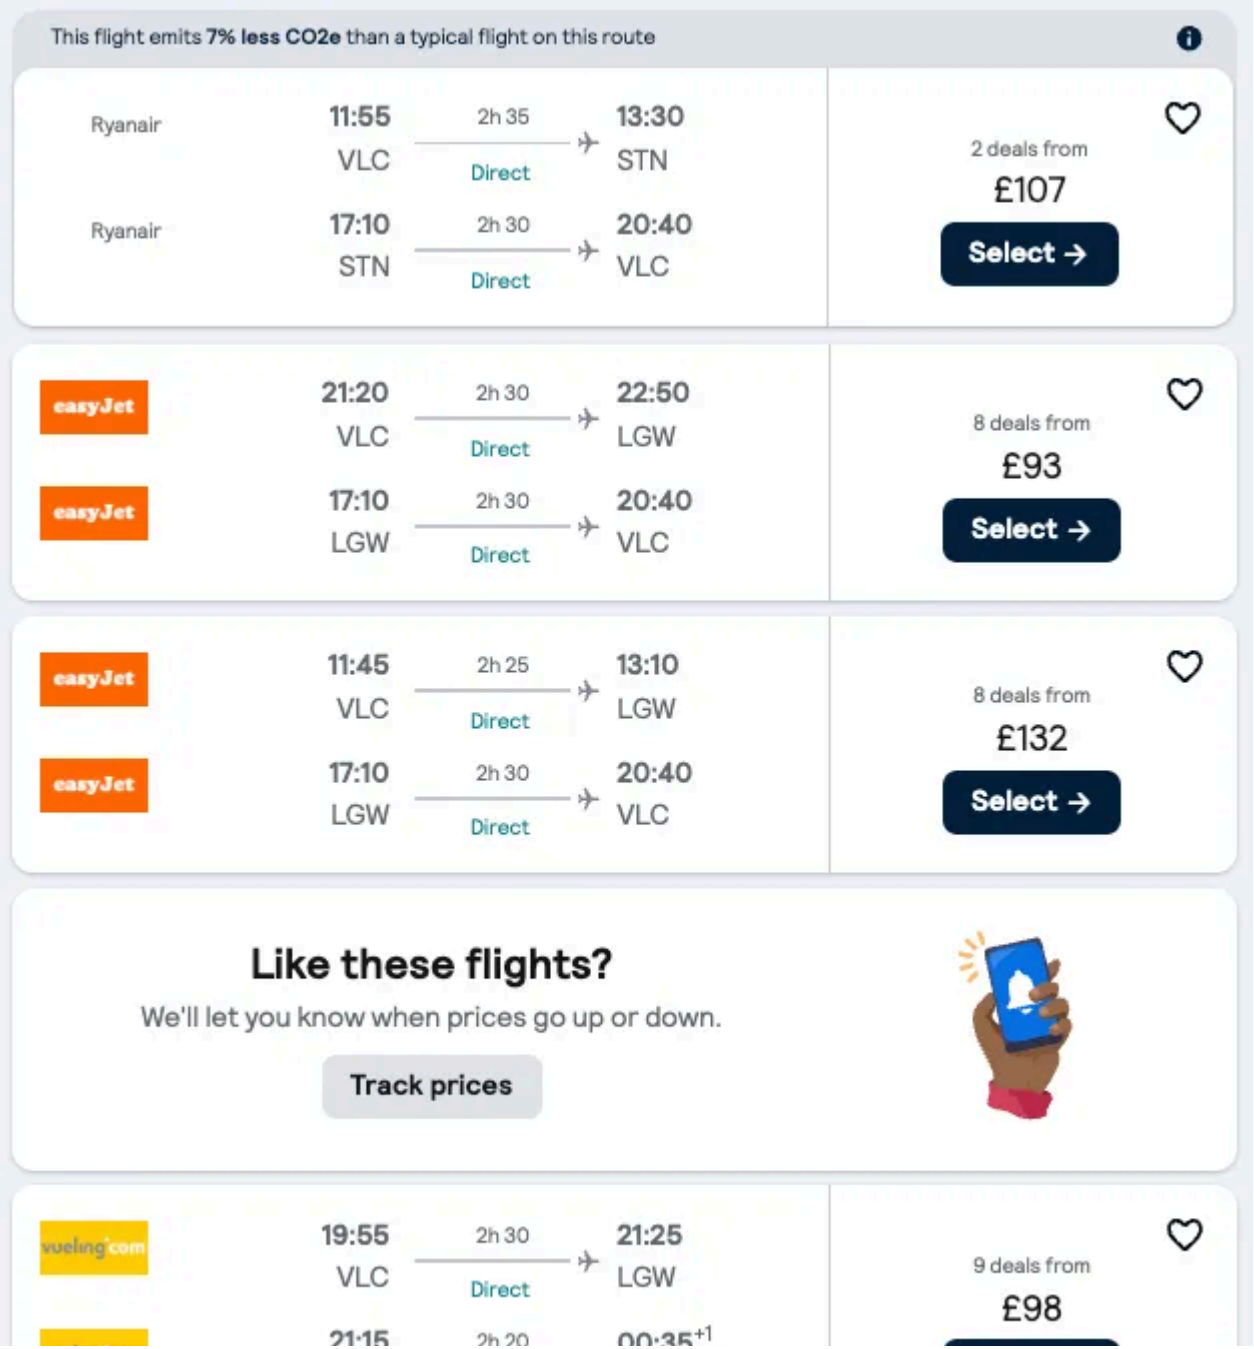
This flight emits 7% less CO2e than a typical flight on this route
Ryanair
11:55
VLC
2h 35
Direct
13:30
STN
Ryanair
17:10
STN
2h 30
Direct
20:40
VLC
2 deals from
£107
Select →
easyJet
21:20
VLC
2h 30
Direct
22:50
LGW
easyJet
17:10
LGW
2h 30
Direct
20:40
VLC
8 deals from
£93
Select →
easyJet
11:45
VLC
2h 25
Direct
13:10
LGW
easyJet
17:10
LGW
2h 30
Direct
20:40
VLC
8 deals from
£132
Select →
Like these flights?
We'll let you know when prices go up or down.
Track prices
vueling.com
19:55
VLC
2h 30
Direct
21:25
LGW
9 deals from
£98

A screenshot of how Skyscanner presents flights options, and how these have to be ranked to help travellers find the best options at the top.

They successfully deployed an experiment that increased redirect rates by small percentage, but leading to an additional £1M in annualised revenue.

A million dollars. That is real, measurable impact.

So, why wasn't this a clear-cut recognition case?

The issue wasn't the outcome. It was how the result was achieved.

Instead of developing a novel ranking approach or strategically enhancing our model, the senior data scientist relied on brute-force optimisation. The approach itself required persistence, patience, and solid execution, but not deep innovation.

And to be clear: **this is not a bad approach.** It's a completely valid way to optimise machine learning models. In fact, for a junior or mid-level data scientist, this would be an outstanding contribution. But for a *senior*? The expectation is **not just to find a better model, it's to think more strategically about how we get there.**

(PS: If you are interested in knowing more about our competency framework, check this article where I go into it in detail)

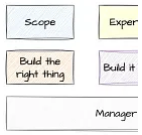

Data Science competency framework: how to define roles and growth

JOSE PARREÑO GARCIA • 8 FEBRUARY 2025

Read full story →

Recognition vs. long-term strategic growth

This is where the distinction between **recognition and career progression** becomes important.

- **Recognition is deserved.** This person delivered real business value. It's entirely reasonable to acknowledge their effort. A stronger bonus or public appreciation are great ways of doing so.
- **But was it a promotion-worthy achievement?** Not necessarily. Strategic, senior-level impact goes beyond just testing variations. It involves pushing boundaries, introducing new frameworks, or fundamentally improving our approach.

If we start promoting based solely on absolute impact, without considering how the impact was achieved, we set the wrong precedent:

- We risk incentivising brute-force, low-leverage solutions instead of encouraging deep problem-solving.
- We signal that optimisation work alone is enough to climb to the next level.

- We fail to differentiate between impact that any capable data scientist could achieve and impact that truly redefines our competitive advantage.

So, the real challenge is to figure out if we are rewarding results, or are we building a culture that encourages the right kind of impact?

Next, let's look at another £1M success story—but this time, one that was fully recognised.

Story #2: Hotel ranking: the other £1M, but this time, it was recognised

> *"Two teams. Two ranking models. Two £1M revenue boosts. So why was this one seen differently?"*

On the surface, the outcomes look identical. The flights ranking experiment and the hotel ranking project both increased revenue by $1M per year. But the effort, complexity, and long-term impact of the two could not have been more different.

- They had to design and implement a brand-new ranking system from scratch.
- They worked cross-functionally with engineers, data scientists, and product teams to get it production-ready.
- They integrated machine learning models, heuristics, and operational constraints to ensure accuracy.

The heuristics they initially used may not have been more complex than the brute-force approach used in flights ranking. But the difference is they weren't just iterating on a model; they were creating the entire infrastructure that powered ranking.

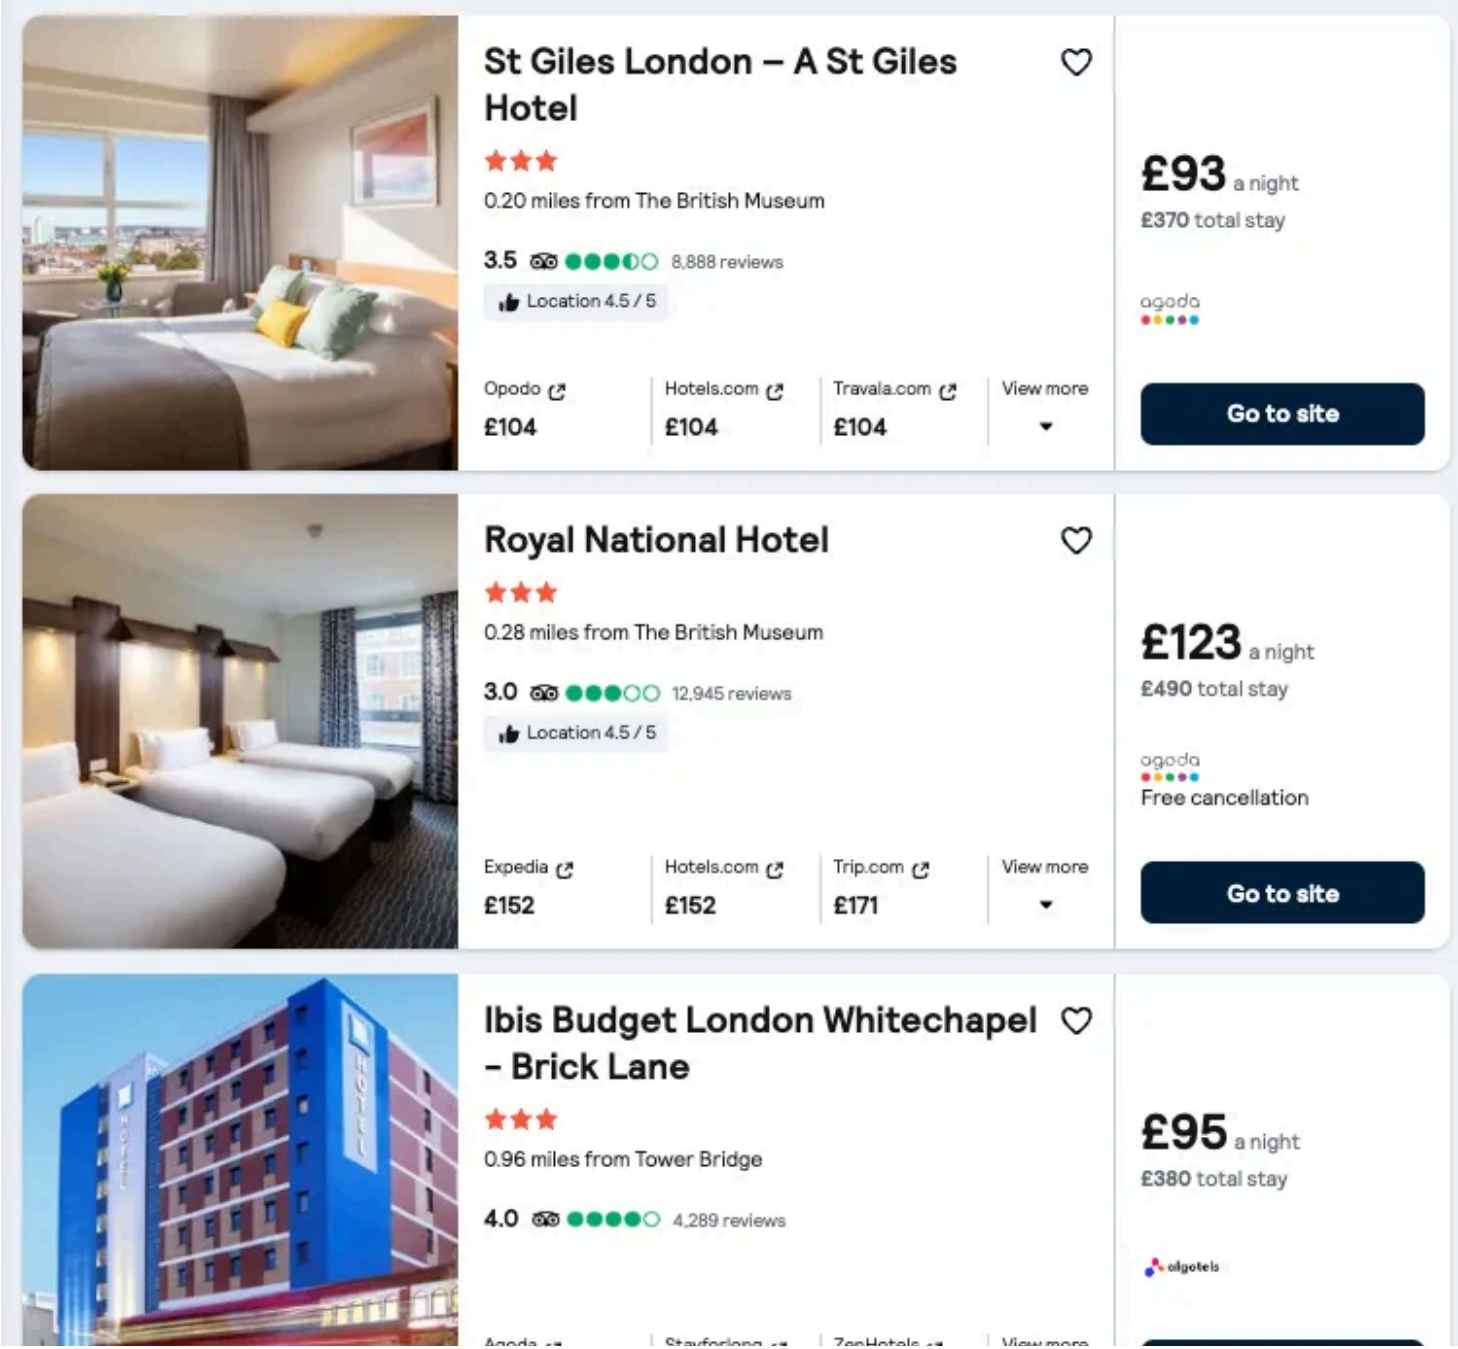
St Giles London – A St Giles Hotel
0.20 miles from The British Museum
3.5 8,888 reviews
Location 4.5 / 5
Opodo £104
Hotels.com £104
Travala.com £104
View more
£93 a night
£370 total stay
agoda
Go to site
Royal National Hotel
0.28 miles from The British Museum
3.0 12,945 reviews
Location 4.5 / 5
Expedia £152
Hotels.com £152
Trip.com £171
View more
£123 a night
£490 total stay
agoda
Free cancellation
Go to site
Ibis Budget London Whitechapel – Brick Lane
0.96 miles from Tower Bridge
4.0 4,289 reviews
£95 a night
£380 total stay
algotels

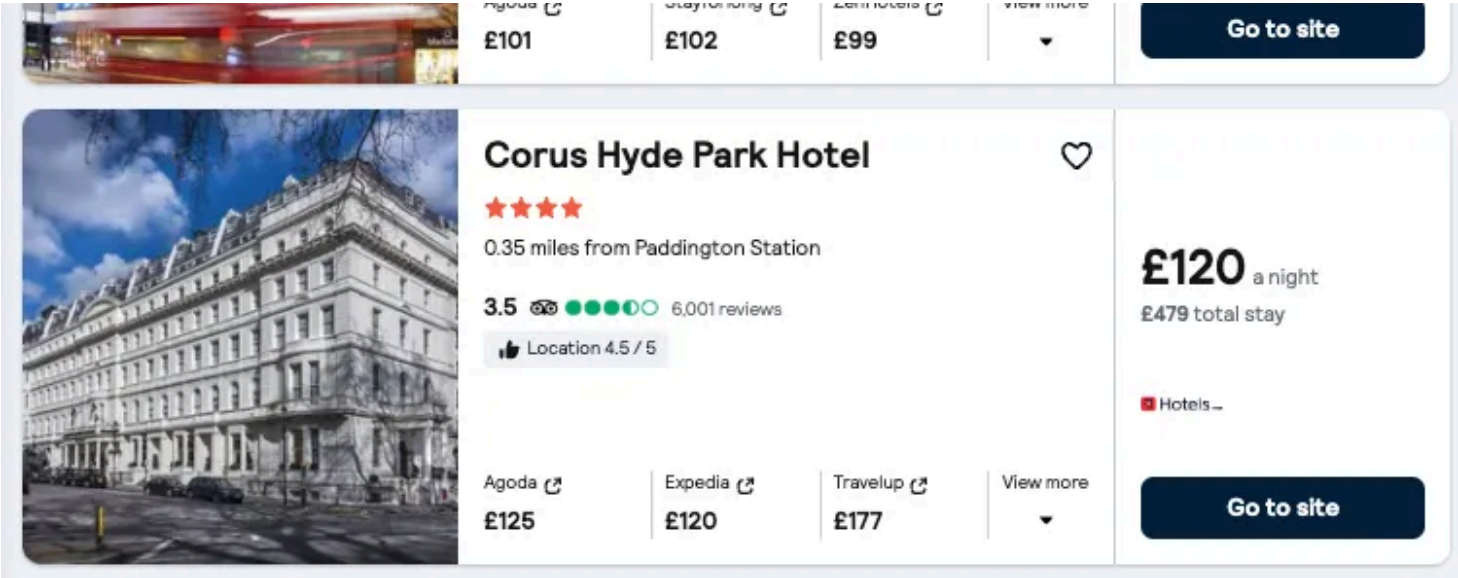

Similarly for flights, hotels also need to be ranked to help travellers find a great place to stay in.

Same financial impact, but a different kind of work

So why did this team receive more recognition than the flights ranking team? Because absolute impact isn't the only factor. **Context matters.**

- **Flights Ranking was an optimisation.** They tuned parameters and ran brute-force experiments to squeeze out incremental gains.
- **Hotel Ranking was an architectural rebuild.** They created an entirely new framework that could support further innovation.

When evaluating impact, we need to ask:

1. **Is this solving a high-leverage problem?** The hotel ranking system improved not just one model but the entire infrastructure, unlocking future enhancements beyond this single £1M win.
2. **Does this scale beyond the immediate project?** The new system can support future features and ML improvements, creating a compounding effect.
3. **Is this work foundational?** Unlike a parameter tweak, this was a major system investment that elevated our ranking capabilities across the company.

In this story, the key element is that we ensured that we recognised and rewarded work that raised the bar.

Next, let's shift the conversation from financial impact to another flawed metric: recognition-driven measurement.

Story #3: Enablement teams: The problem with visibility-based recognition

> *"If a team gets 10x more public praise, does that mean they are 10x more impactful?"*

So far, we have examined impact through the lens of measurable business value. But what happens when impact is measured by recognition rather than actual contribution?

This is where visibility bias comes into play.

In our data science discipline, we have a centralised enablement team responsible for maintaining and improving our A/B testing statistical engine.

As part of their role, they:

- Provide support for teams across the company running experiments.
- Host “green flag” Q&A sessions to help teams interpret results.
- Ensure that our internal experimentation framework is used correctly.

Because of their cross-functional nature, they interact with far more teams than a typical data science squad. And since they are often solving urgent, visible problems, they receive a great amount of recognition.

We have an internal “high-5” system where employees can send thank-you notes and public appreciation to colleagues. It’s a great tool for celebrating contributions, but if we were to measure impact by sheer number of high-5s, the enablement team would outperform every other team by an order of magnitude.

- Does that mean they are 10x more impactful than a team working on core ranking models?
- Does that mean they should be prioritised for promotions over others who work on complex, but less visible, problems?

Of course not. But this is exactly what happens when we conflate recognition with impact.

So, how do we know if enablement teams are overperforming?

This doesn't mean the enablement team isn't valuable. Let me be clear: they are. Their role is to amplify the effectiveness of others. But measuring their impact solely by recognition would be like judging an engineer's effectiveness by the number of Slack messages they send.

Instead, a good way of looking at contributions from enablement teams would be:

- Look at the multiplier effect: how much do they improve the quality of experimentation across the company?
- What tooling did they ship that helped teams move forward in new ways?
- How many green flag duties do they have to cover? Paradoxically, if the platform was perfect, the team wouldn't get too many urgent requests.

Next, let's look at another case where traditional impact measurement falls short: slow-moving metrics that take years to show their full value.

Story #4: Slow metrics, where long-term impact is hard to measure

> *"If a team spends three years improving a system, but the results come in tiny increments, does that make their work any less valuable?"*

Some of the most important work in a company happens slowly, bit by bit, over months or even years. If we are not careful, we risk undervaluing it simply because it doesn't fit neatly into quarterly or annual reporting cycles.

The flight price calendar: A 3-year effort with an astonishing result

One of the most complex challenges in flight meta-search is accurately estimating flight prices over time. Imagine a traveler looking at our calendar view, where they can see estimated prices for flights across an entire month or even a full year.

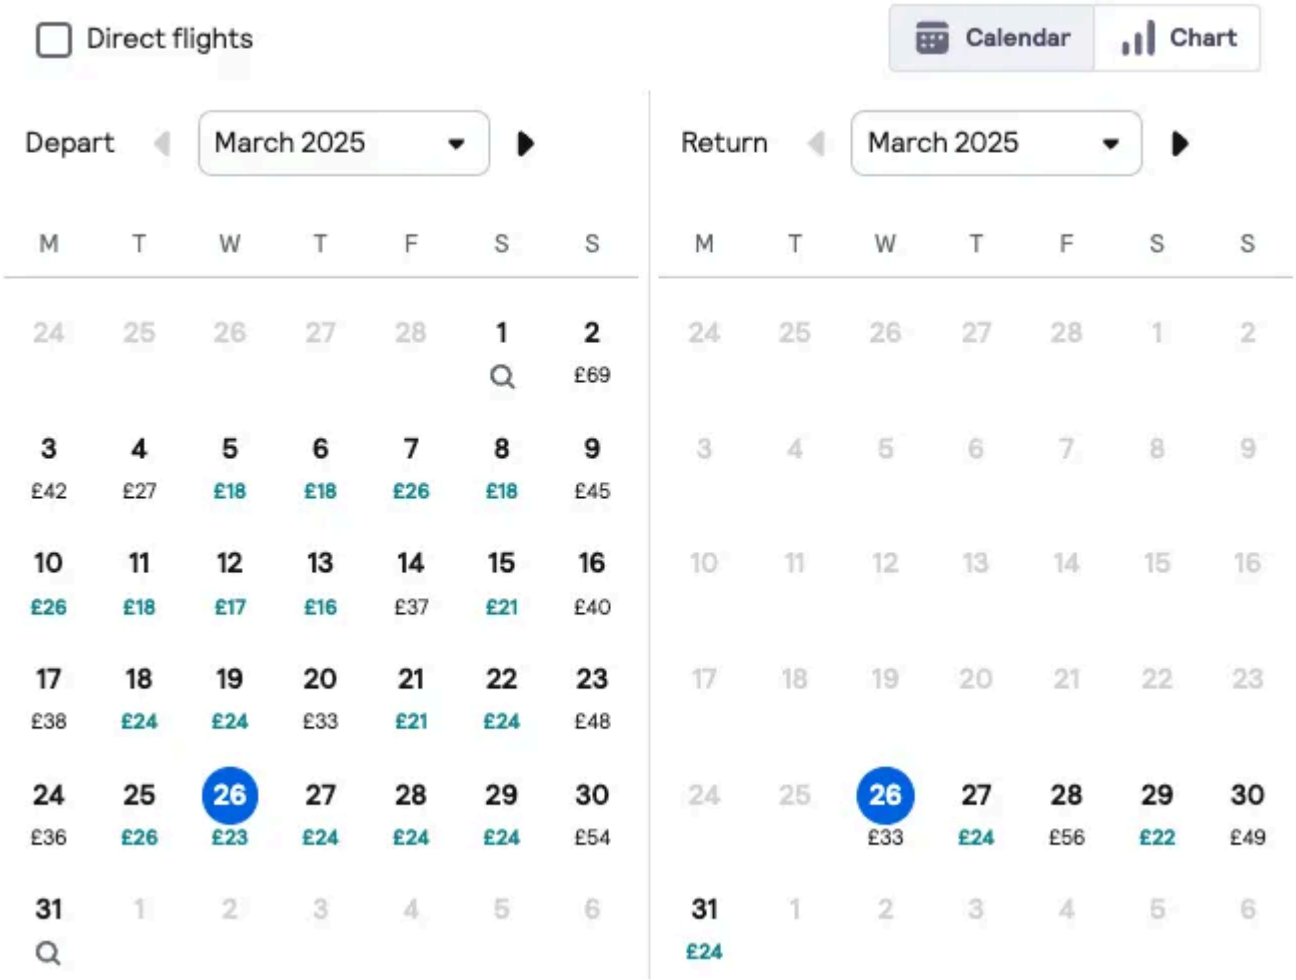

Showing prices in a calendar is actually *really* difficult, even more to be able to fill out the full date grid!

Seems straightforward, right? I can tell you, it is not.

- **Flight prices change constantly**, sometimes multiple times per day.

- **We don't always have real-time data** for every date, meaning we need to estimate missing prices.
- **Users expect accuracy**. If they see a flight for £60 on the calendar but land on a page where it's suddenly £80, they lose trust in our platform.

For years, our coverage accuracy metric hovered at 40%. Because of that, our calendar view was super sparse, as we didn't want to fill it with prices where we knew that 60% of the times they were wrong.

A dedicated team spent 3 years improving this system.

- They developed new estimation techniques to increase accuracy.
- They improved how we pull and refresh pricing data at scale.
- They fine-tuned coverage algorithms to ensure fewer gaps in the calendar.

And after 3 years of continuous work, they raised the **accuracy metric from 40% to 80%.**

Why this work was at risk of being overlooked

This improvement happened bit by bit, quarter by quarter. No single release made a giant, eye-catching difference overnight. But when you zoom out, the impact is undeniable:

- 80% accuracy means a dramatically better user experience.
- Users trust the calendar view more, increasing engagement and bookings.
- The foundation is now in place for future improvements, making our pricing system far more scalable.

If we are not intentional, slow-moving improvements like this get ignored because:

1. **We expect impact to be immediate.** Quarterly results dominate decision-making, making multi-year investments seem "less urgent."
2. **It's hard to tie improvements directly to revenue.** Unlike an A/B test that increases conversions by 1%, this work supports many different downstream impacts.
3. **Incremental gains are harder to "sell" internally.** A team that says, *"We improved accuracy from 40% to 43% this quarter"* sounds less impressive than one that says, *"We ran an experiment that made $500K."*

So, next time, ask yourself if you are looking at the full picture, not just quarterly snapshots? If you are recognising compounding improvements, even if they take years? Or, if you are valuing long-term thinking as much as short-term wins?

Finally, let's look at an even harder-to-measure form of impact: The Amplifier Effect.

Story #5: The amplifier effect

> *"What if the most valuable work doesn't move a single metric—but makes everything better?"*

We have explored cases where impact was measured incorrectly, undervalued, or too slow to be appreciated. But what about work that doesn't have a direct metric at all? Work that multiplies the effectiveness of others, but is never attached to a single person's performance review?

A Data Scientist who transformed engineering standards

One of the most impactful people on my team wasn't a senior engineer or a principal scientist. She was a mid-level data scientist who saw a problem, took ownership, and quietly reshaped how we worked.

- She noticed that our engineering standards for Git repos, data storage, and other core infrastructure were inconsistent.
- Instead of waiting for someone else to fix it, she took the initiative to improve them.
- She reached out to senior engineers for guidance, learning how to implement best practices properly.
- She built a new, standardised approach that the team agreed was the right way forward.

The improvements weren't tied to a specific revenue goal. They weren't even officially on the roadmap. But, they made everything smoother, faster, and more maintainable.

Then, came the most important part. Instead of moving on to another project, she focused on spreading what she learned.

- She trained 3 other teammates on the new framework.
- She ensured no single person was a bottleneck for this knowledge.
- She created a standard approach that future projects could build on.

The problem with traditional impact measurement

If you judged her performance by standard metrics, her work didn't show up anywhere.

- She didn't build a new model.
- She didn't deliver a direct revenue increase.
- She didn't improve an existing KPI.

But what she did do was de-risk future projects, improve team efficiency, and level up the entire organisation. And as part of our competency framework, she nailed the "expertise" and "build it right" sections for her level. Thus, was accordingly rewarded with a great bonus and more challenges that would set her up to a promotion path.

Rethinking how we measure impact

> *"Not all impact is measurable. And not all measurable impact is equal."*

I hope you have enjoyed my 5 stories that challenged me personally to change the way I was defining impact. To summarise them:

- **Impact isn't just about revenue.** Two $1M projects can have vastly different strategic value.
- **Visibility doesn't equal value.** The loudest contributions aren't always the most important.
- **Slow-moving progress matters.** Some of the most critical work takes years to show results.
- **Amplifiers shape the future.** The best contributors don't just create impact—they enable it at scale.

If we only reward what is easiest to measure, we risk deprioritising foundational work, discouraging long-term thinking, and failing to recognise those who make the entire system better.

Now, I want to hear from you!

🧠 How do *you* recognise great performance?

- What are the signs you look for — especially when impact isn't immediate or easily measured?
- Have you ever seen someone's work overlooked because it didn't show up in a roadmap or OKR?

- Or the opposite — work celebrated *loudly* but with questionable long-term value?
- And if you're a manager: how do you balance consistency in your process with flexibility for context?

I'd love to hear how you think about performance, visibility, long-term impact, and recognising the "quiet stars" on your team.

Drop your thoughts in the comments, especially if you have ever questioned what your performance process is really optimising for.

And finally, another big thanks to Rafa Páez for collaborating with me on this leadership topic!

Further reading

If you are interested in more content, here is an article capturing all my written blogs!

All my articles, one convenient place

JOSE PARREÑO GARCIA • 15 OCTOBER 2024

Read full story →

3 Likes • 1 Restack

A guest post by

Rafa Páez

Engineering Leader and AI Enthusiast. Sharing insights from 20+ years of experience. I've scaled platforms and teams at unicorn startups in remote environments. I enjoy helping others speed up their careers.

Subscribe to Rafa

Discussion about this post

Comments Restacks

Write a comment...

© 2026 Jose Parreño Garcia · Privacy · Terms · Collection notice
Substack is the home for great culture